

# Camunda 8 Helm Chart

License

Apache 2.0

 Test - Unit

failing

 Camunda Platform

10.0.4

Please also refer to the documentation on how to use Helm charts.

- Architecture
- Requirements
- Dependencies
- Versioning
- Installation
  - Local Kubernetes
  - OpenShift
- Backporting
- Uninstalling Charts
- Configuration
- Notes on Configuration
  - Web Modeler
  - Elasticsearch
  - Keycloak
- Development
- Releasing the Charts
- Parameters
  - Global parameters
  - Console Parameters
  - Zeebe Parameters
  - ZeebeGateway Parameters
  - Operate Parameters
  - Tasklist Parameters
  - Optimize Parameters
  - Identity Parameters
  - Identity - PostgreSQL Parameters
  - Identity - Keycloak Parameters
  - WebModeler Parameters
  - WebModeler - RestAPI Parameters
  - WebModeler - WebApp Parameters
  - WebModeler - WebSockets Parameters
  - WebModeler - PostgreSQL Parameters
  - Connectors Parameters
  - Elasticsearch Parameters
  - Prometheus Parameters

## Architecture

## Requirements

- Helm >= 3.9.x
- Kubernetes >= 1.20+
- Minimum cluster requirements include the following to run this chart with default settings.
  - All of these settings are configurable.
  - Three Kubernetes nodes to respect the default “hard” affinity settings
  - 2GB of RAM for the JVM heap

## Dependencies

Camunda 8 Helm chart is an umbrella chart for different components. Some are internal (sub-charts), and some are external (third-party). The dependency management is fully automated and managed by Helm itself; however, it’s good to understand the dependency structure. This third-party dependency is reflected in the Helm chart as follows:

```
camunda-platform
├─ _elasticsearch
├─ _identity
│  └─ _keycloak
│     └─ _postgresql
├─ _optimize
├─ _operate
└─ _tasklist
```

```
|_ zeebe
|_ postgresql
```

[!NOTE] Please note that the Connectors and Web Modeler components are part of the main chart and not implemented as sub-charts.

For example, Camunda Identity utilizes Keycloak and allows you to manage users, roles, and permissions for Camunda 8 components.

- Keycloak is a dependency for Camunda Identity, and PostgreSQL is a dependency for Keycloak.
- Elasticsearch is a dependency for the Camunda chart, which is used in Zeebe, Operate, Tasklist, and Optimize.
- PostgreSQL is an optional dependency for the Camunda chart and is used by Web Modeler.

The values for the dependencies Keycloak and PostgreSQL can be set in the same hierarchy:

```
identity:
  [identity values]
keycloak:
  [keycloak values]
postgresql:
  [postgresql values]
postgresql:
  [postgresql values]
```

## Versioning

After the 8.4 release (January 2024), the Camunda Helm chart version is **decoupled** from the version of the application (e.g., the chart version is 9.0.0 and the application version is 8.4.x).

Before the 8.4 release, the Camunda Helm chart version was **coupled** with the applications version (e.g., chart version is 8.3.x and applications version is 8.3.x).

For more details, check out the full version matrix.

## Installation

The first command adds the official Camunda Helm charts repo, and the second installs the Camunda chart to your current Kubernetes context.

```
helm repo add camunda https://helm.camunda.io
helm install camunda-platform camunda/camunda-platform
```

Although the Camunda 8 Helm chart gets the latest version of Camunda 8 applications, the version is still possible to diverge slightly between the chart and the apps (more details about that can be found in versioning).

To have the latest version of the chart and apps at any time, install the chart as follows:

```
helm install camunda-platform camunda/camunda-platform \
  --values https://helm.camunda.io/camunda-platform/values/values-latest.yaml
```

For the previous version, you can get the latest applications patch version using our backporting mechanism.

### Local Kubernetes

We recommend using Helm on KIND for local environments, as the Helm configurations are battle-tested and much closer to production systems.

For more details, follow the Camunda 8 local Kubernetes cluster guide.

### OpenShift

Check out OpenShift Support to get started with deploying the charts on Red Hat OpenShift.

## Backporting

Our Helm chart is highly customizable and constantly evolving. Hence, currently, we backport the older charts by providing an extra value file per version. That covers most backporting cases, like updating the application’s image tags to the latest patch version, setting env var, etc.

To install a previous chart version with the latest app patch image tags for that version, use the values file for the minor release. For example (the values file could also be downloaded):

```
helm install camunda-platform camunda/camunda-platform --version 8.1 \
  --values https://helm.camunda.io/camunda-platform/values/values-v8.1.yaml
```

## Uninstalling Charts

You can remove these charts by running:

```
helm uninstall camunda
```

[!NOTE]

Notice that all the Services and Pods will be deleted, but not the PersistentVolumeClaims (PVC) which are used to hold the storage for the data generated by the cluster and Elasticsearch.

To free up the storage, you need to delete all the PVCs manually.

First, view the PVCs:

```
kubect1 get pvc -l app.kubernetes.io/instance=camunda
kubect1 get pvc -l release=camunda
```

Then delete the ones that you don't want to keep:

```
kubect1 delete pvc -l app.kubernetes.io/instance=camunda
kubect1 delete pvc -l release=camunda
```

Or you can delete the related Kubernetes namespace, which contains all PVCs.

Configuration

The following sections contain the configuration values for the chart and each sub-chart. All of them can be overwritten via a separate values.yaml file.

Check out the default values.yaml file, which contains the same content and documentation.

[!NOTE]

For more details about deploying Camunda 8 on Kubernetes, please visit the Helm/Kubernetes installation instructions docs.

Notes on Configuration

Web Modeler

[!NOTE]

Web Modeler Self-Managed is available to Camunda enterprise customers only.

**Docker registry** The Docker images for Web Modeler are available in a private registry. Enterprise customers either already have credentials to this registry, or they can request access to this registry through their CSM contact at Camunda. To enable Kubernetes to pull the images from Camunda's registry, you'll need to:

- create an image pull secret using the provided credentials
- configure the Web Modeler pods to use the secret:

```
webModeler:
  image:
    pullSecrets:
      - name: <SECRET_NAME>
```

**Database** Web Modeler requires a PostgreSQL database to store the data. You can either:

- Deploy a PostgreSQL instance as part of the Helm release by setting postgresql.enabled to true (which will enable the postgresql chart dependency).
- Configure a connection to an (existing) external database by setting postgresql.enabled to false and providing the values under restapi.externalDatabase.

**SMTP server** Web Modeler requires an SMTP server to send (notification) emails to users. The SMTP connection can be configured with the values under restapi.mail.

**Updating Environment Variables** When configuring the env options in the settings listed above, the environment variables you specify in values.yaml may show up twice when running kubect1 describe deployment <deployment>. However, the environment variable is specified in values.yaml will have precedence when the pod actually runs. To verify this, you can check the output from the following command:

```
kubect1 exec pod/<podName> -- env
```

**Outbound Connectors** To learn more about outbound connectors, visit related documentation article.

**Inbound Connectors** To learn more about inbound connectors, visit related documentation article.

**Using Connector Secrets** Connector secrets are generally configured via environment variables.

You can set them via values.yaml, or command line. For example, if you need to set a Slack token, you should configure the following:

```
connectors:
  env:
    - name: SLACK_TOKEN
      value: <your actual token value>
```

After that, a Modeler user can set in their BPMN diagram a value secrets.SLACK\_TOKEN without ever knowing the actual token.

Visit using secrets in manual installation to learn more.

Elasticsearch

Camunda 8 Helm chart has a dependency on the Elasticsearch 8 Helm Chart. All variables related to Elasticsearch can be set under `elasticsearch`.

[!NOTE]

The default setup of the Elasticsearch 8 part of Camunda 8 uses nodes that have all roles (master, data, coordinating, and ingest). For high-demand deployments, it’s recommended to deploy the Elasticsearch master-eligible nodes as master-only nodes.

Section	Parameter	Description	Default
elasticsearch	enabled	If true, enables Elasticsearch deployment as part of the Camunda Helm chart	true

Example:

```
elasticsearch:
  enabled: true
  image:
    tag: <YOUR_VERSION_HERE>
```

**Elasticsearch Retention** Since moving to Elasticsearch 8, Curator is deprecated in favor of Manage the index lifecycle (ILM). Hence, each component in Camunda 8 controls its Elasticsearch index retention.

Keycloak

When Camunda 8 Identity component is enabled by default, and it depends on Bitnami Keycloak chart. Since Keycloak is a dependency for Identity, all variables related to Keycloak can be found in `bitnami/keycloak/values.yaml` and can be set under `identity.keycloak`.

Section	Parameter	Description	Default
identity.keycloak	enabled	If true, enables Keycloak chart deployment as part of the Camunda Helm chart	true

Example:

```
identity:
  keycloak:
    enabled: true
```

**Keycloak Theme** Camunda provides a custom theme for the login page used in all apps. The theme is copied from the Identity image.

The theme is added to Keycloak by default, however, since Helm v3 (the latest checked 3.10.x) doesn’t merge lists with custom values files, then you will need to add this to your own values file if you override any of `extraVolumes`, `initContainers`, or `extraVolumeMounts`.

```
identity:
  keycloak:
    extraVolumes:
      - name: camunda-theme
        emptyDir:
          sizeLimit: 10Mi
    initContainers:
      - name: copy-camunda-theme
        image: >-
          {{- $identityImageParams := (dict "base" .Values.global "overlay" .Values.global.identity) -}}
          {{- include "camundaPlatform.imageByParams" $identityImageParams }}
        imagePullPolicy: "{{ .Values.global.image.pullPolicy }}"
        command: ["sh", "-c", "cp -a /app/keycloak-theme/* /mnt"]
        volumeMounts:
          - name: camunda-theme
            mountPath: /mnt
    extraVolumeMounts:
      - name: camunda-theme
        mountPath: /opt/bitnami/keycloak/themes/identity
```

Development

For development purposes, you might want to deploy and test the charts without creating a new helm chart release. To do this you can run the following:

```
helm install camunda --atomic --debug ./charts/camunda-platform
```

- --atomic if set, the installation process deletes the installation on failure. The --wait flag will be set automatically if --atomic is used
- --debug enable verbose output

To generate the resources/manifests without really installing them, you can use:

- --dry-run simulate an install

If you see errors like:

Error: found in Chart.yaml, but missing in charts/ directory: elasticsearch

Then you need to download the dependencies first.

Run the following to add resolve the dependencies:

```
make helm.repos-add
```

After this, you can run: make helm.dependency-update, which will update and download the dependencies for all charts.

The execution should look like this:

```
$ make helm.dependency-update
helm dependency update charts/camunda-platform
Hang tight while we grab the latest from your chart repositories...
...Successfully got an update from the "camunda-platform" chart repository
...Successfully got an update from the "bitnami" chart repository
Update Complete. Happy Helming!
Saving 6 charts
Dependency zeebe did not declare a repository. Assuming it exists in the charts directory
Dependency zeebe-gateway did not declare a repository. Assuming it exists in the charts directory
Dependency operate did not declare a repository. Assuming it exists in the charts directory
Dependency tasklist did not declare a repository. Assuming it exists in the charts directory
Dependency identity did not declare a repository. Assuming it exists in the charts directory
Deleting outdated charts
helm dependency update charts/camunda-platform/charts/identity
Hang tight while we grab the latest from your chart repositories...
...Successfully got an update from the "camunda-platform" chart repository
...Successfully got an update from the "bitnami" chart repository
Update Complete. Happy Helming!
Saving 2 charts
Downloading keycloak from repo https://charts.bitnami.com/bitnami
Downloading common from repo https://charts.bitnami.com/bitnami
```

Releasing the Charts

Please see the corresponding release guide to find out how to release the chart.

Parameters

Global parameters

global		
global.multitenancy		
global.multitenancy.enabled	if true, then enable multitenancy in all applicable components.	false
global.createReleaseInfo	Create config that will be used in Camunda Console.	true
global.annotations	Annotations can be used to define common annotations, which should be applied to all deployments	{}
global.labels.app	Name of the application	camunda-platform
global.image.registry	Can be used to set container image registry.	""
global.image.tag	defines the tag / version which should be used in the most of the apps.	8.5.0
global.image.pullPolicy	defines the image pull policy which should be used	IfNotPresent
	https://kubernetes.io/docs/concepts/containers/images/#image-pull-policy	
global.image.pullSecrets	can be used to configure image pull secrets	[]
	https://kubernetes.io/docs/concepts/containers/images/#specifying-imagepullsecrets-on-a-pod	

global.ingress		
global.ingress.enabled	if true, an ingress resource is deployed. Only useful if an ingress controller is available, like Ingress-NGINX.	false
global.ingress.className	Ingress.className defines the class or configuration of ingress which should be used by the controller	nginx
global.ingress.annotations	defines the ingress related annotations, consumed mostly by the ingress controller	{}
global.ingress.host	If not specified the rules applies to all inbound http traffic, if specified the rule applies to that host.	""
global.ingress.pathType	can be used to define the Ingress path type. <a href="https://kubernetes.io/docs/concepts/services-networking/ingress/#path-types">https://kubernetes.io/docs/concepts/services-networking/ingress/#path-types</a>	Prefix
global.ingress.tls	configuration for tls on the ingress resource <a href="https://kubernetes.io/docs/concepts/services-networking/ingress/#tls">https://kubernetes.io/docs/concepts/services-networking/ingress/#tls</a>	
global.ingress.tls.enabled	if true, then tls is configured on the ingress resource. If enabled the Ingress.host need to be defined.	false
global.ingress.tls.secretName	defines the secret name which contains the TLS private key and certificate	camunda-platform
global.elasticsearch		
global.elasticsearch.enabled	if true, enables elasticsearch for all components	true
global.elasticsearch.external	if true, tries to connect to an external elasticsearch	false
global.elasticsearch.tls		
global.elasticsearch.tls.enabled	enable tls for external elasticsearch	false
global.elasticsearch.tls.existingSecret	provide an already existing tls secret for connecting to external elasticsearch	nil
global.elasticsearch.auth		
global.elasticsearch.auth.username	the username for external elasticsearch	nil
global.elasticsearch.auth.password	the password for external elasticsearch	nil
global.elasticsearch.auth.existingSecret	you can provide an existing secret for the external elasticsearch password	nil
global.elasticsearch.auth.existingSecretKey	you can provide an existing secret key for the external elasticsearch password	nil
global.elasticsearch.disableExporter	DEPRECATED: this value is not needed anymore. Use global.elasticsearch.enabled	false
global.elasticsearch.url	Configuration to configure elasticsearch url	
global.elasticsearch.url.protocol	defines the elasticsearch access protocol.	http
global.elasticsearch.url.host	Elasticsearch.host defines the elasticsearch host, ideally the service name inside the namespace	{{ .Release.Name }}-elasticsearch
global.elasticsearch.url.port	Elasticsearch.port defines the elasticsearch port, under which elasticsearch can be accessed	9200
global.elasticsearch.clusterName	Elasticsearch.clusterName defines the cluster name which is used by Elasticsearch	elasticsearch
global.elasticsearch.prefix	Elasticsearch.prefix defines the prefix which is used by the Zeebe Elasticsearch Exporter to create Elasticsearch indexes	zeebe-record
global.opensearch		
global.opensearch.enabled	enabled external opensearch	false
global.opensearch.aws.enabled	Enabling AWS IRSA	false
global.opensearch.tls		
global.opensearch.tls.enabled	enable tls for external opensearch	false
global.opensearch.tls.existingSecret	provide an already existing tls secret for connecting to external opensearch	nil
global.opensearch.auth		
global.opensearch.auth.username	the username for external opensearch	nil
global.opensearch.auth.password	the password for external opensearch	nil
global.opensearch.auth.existingSecret	you can provide an existing secret for the external opensearch password	nil
global.opensearch.auth.existingSecretKey	you can provide an existing secret key for the external opensearch password	nil
global.opensearch.url	Configuration to configure opensearch url	
global.opensearch.url.protocol	defines the external opensearch access protocol	https
global.opensearch.url.host	defines the external opensearch host, ideally the service name inside the namespace	nil
global.opensearch.url.port	defines the external opensearch port, under which opensearch can be accessed	443
global.zeebeClusterName	ZeebeClusterName defines the cluster name for the Zeebe cluster. All Zeebe pods get this prefix in their name and the brokers uses that as cluster name.	{{ .Release.Name }}-zeebe
global.identity.keycloak.internal	It's useful for using existing Keycloak in another namespace with and access it with the combined Ingress.	false
global.identity.keycloak.url	can be used incorporate with “identityKeycloak.enabled: false” to use your own Keycloak instead of the one comes with Camunda Helm chart.	{}
global.identity.keycloak.contextPath	In Keycloak v16.x.x it's hard-coded as ‘/auth’, but in v19.x.x it's ‘/’.	/auth
global.identity.keycloak.realm	defines Keycloak realm path used for Camunda.	/realms/camunda-platform
global.identity.keycloak.auth	same as “identityKeycloak.auth” but it's used for existing Keycloak.	{}
global.identity.auth	configuration, to configure identity authentication setup	
global.identity.auth.enabled	if true, enables the identity authentication otherwise basic-auth will be used on all services.	true
global.identity.auth.issuer	defines the issuer name, which is used by the services to validate the JWT tokens.	""

global.identity.auth.issuerBackendUrl	defines the issuer backend URL, which is used by the services to validate the JWT tokens in a container to container context.	""
global.identity.auth.tokenUrl	defines the token URL, which is used by the services to request JWT tokens.	""
global.identity.auth.jwksUrl	defines the JWKS URL, which is used by the services to validate the JWT tokens.	""
global.identity.auth.type	defines the type of authentication which should be used. Defaults to Keycloak	KEYCLOAK
global.identity.auth.publicIssuerUrl	Can be overwritten if ingress is in use and an external IP is available.	http://localhost:18080/auth/realms/camunda-platform
global.identity.auth.connectors	configuration to configure Connectors authentication specifics on global level, which can be accessed by other sub-charts	
global.identity.auth.connectors.clientId	defines the client id, which is used by Connectors in authentication flows.	connectors
global.identity.auth.connectors.existingSecret	can be used to use an own existing secret. If not set a random secret is generated.	""
global.identity.auth.identity	configuration to configure Identity authentication specifics on global level, which can be accessed by other sub-charts	
global.identity.auth.identity.clientId	defines the client id, which is used by Identity in authentication flows.	identity
global.identity.auth.identity.audience	defines the audience, which is used by Identity.	camunda-identity-resource-server
global.identity.auth.identity.existingSecret	can be used to reference an existing secret. If not set, a random secret is generated.	nil
global.identity.auth.identity.redirectUrl	defines the redirect URL, which is used by the auth platform to access Identity.	http://localhost:8085
global.identity.auth.identity.initialClaimName	defines the initial claim name, which is used by Identity to configure initial mapping rules,	oid
global.identity.auth.identity.initialClaimValue	defines the initial claim value, which is used by Identity to configure initial mapping rules.	nil
global.identity.auth.operate	configuration to configure Operate authentication specifics on global level, which can be accessed by other sub-charts	
global.identity.auth.operate.clientId	defines the client id, which is used by Operate in authentication flows.	operate
global.identity.auth.operate.audience	defines the audience, which is used by Operate.	operate-api
global.identity.auth.operate.existingSecret	can be used to reference an existing secret. If not set, a random secret is generated.	nil
global.identity.auth.operate.redirectUrl	defines the redirect URL, which is used by Keycloak to access Operate.	http://localhost:8081
global.identity.auth.tasklist	configuration to configure Tasklist authentication specifics on global level, which can be accessed by other sub-charts	
global.identity.auth.tasklist.clientId	defines the client id, which is used by Tasklist in authentication flows.	tasklist
global.identity.auth.tasklist.audience	defines the audience, which is used by Tasklist.	tasklist-api
global.identity.auth.tasklist.existingSecret	can be used to use an own existing secret. If not set a random secret is generated.	nil
global.identity.auth.tasklist.redirectUrl	defines the root (or redirect) URL, which is used by Keycloak to access Tasklist.	http://localhost:8082
global.identity.auth.optimize	configuration to configure Optimize authentication specifics on global level, which can be accessed by other sub-charts	
global.identity.auth.optimize.clientId	defines the client id, which is used by Optimize in authentication flows.	optimize
global.identity.auth.optimize.audience	defines the audience, which is used by Optimize.	optimize-api
global.identity.auth.optimize.existingSecret	can be used to use an own existing secret. If not set a random secret is generated.	nil
global.identity.auth.optimize.redirectUrl	defines the root (or redirect) URL, which is used by Keycloak to access Optimize.	http://localhost:8083
global.identity.auth.webModeler	configuration to configure WebModeler authentication specifics on global level, which can be accessed by other sub-charts	
global.identity.auth.webModeler.clientId	defines the client id, which is used by WebModeler in authentication flows.	web-modeler
global.identity.auth.webModeler.clientApiAudience	defines the audience which is used by WebModeler's client API.	web-modeler-api
global.identity.auth.webModeler.publicApiAudience	defines the audience which is used by WebModeler's public API.	web-modeler-public-api
global.identity.auth.webModeler.redirectUrl	defines the root URL which is used by Keycloak to access WebModeler.	http://localhost:8084
global.identity.auth.console	configuration to configure Console authentication specifics on global level, which can be accessed by other sub-charts	
global.identity.auth.console.existingSecret	can be used to use an own existing secret. If not set a random secret is generated.	nil
global.identity.auth.console.redirectUrl	defines the root URL which is used by Keycloak to access WebModeler.	http://localhost:8080
global.identity.auth.console.audience	can be used to Console audience in Identity.	console-api
global.identity.auth.zeebe	configuration to configure Zeebe authentication specifics on global level, which can be accessed by other sub-charts	
global.identity.auth.zeebe.clientId	defines the client id, which is used by Zeebe in authentication flows.	zeebe
global.identity.auth.zeebe.existingSecret	can be used to use an own existing secret. If not set a random secret is generated.	""
global.identity.auth.zeebe.audience	defines the audience, which is used by Zeebe.	zeebe-api
global.identity.auth.zeebe.tokenScope	defines the token scope, which is used by Zeebe.	nil

### Console Parameters

Name	Description	Value
console	configuration for the Console.	
console.enabled	if true, the Console deployment and its related resources are deployed via a helm release	false

Name	Description	Value
console.configuration	Configuration passed directly to Console as YAML file. More details on Console official documentations	""
console.image.registry	can be used to set container image registry.	registry.camunda.cloud
console.image.repository	defines which image repository to use	console/console-sm
console.image.tag	can be used to set the Docker image tag for the Console image (overwrites global.image.tag)	8.5.4
console.image.pullSecrets	can be used to configure image pull secrets <a href="https://kubernetes.io/docs/concepts/containers/images/#specifying-imagepullsecrets-on-a-pod">https://kubernetes.io/docs/concepts/containers/images/#specifying-imagepullsecrets-on-a-pod</a>	[]
console.sidecars	can be used to attach extra containers to the console deployment	[]
console.replicas	Number of Console replicas	1
console.contextPath	can be used to make Console web application works on a custom sub-path. This is mainly used to run Camunda web applications under a single domain.	""
console.initContainers	can be used to set up extra init containers for the application Pod	[]
console.podAnnotations	can be used to define extra Console pod annotations	{}
console.podLabels	can be used to define extra Console pod labels	{}
console.logging	configuration for the Console logging. This template will be directly included in the Operate configuration YAML file	{}
console.service.annotations	can be used to define annotations, which will be applied to the Console service	{}
console.service.type	defines the type of the service <a href="https://kubernetes.io/docs/concepts/services-networking/service/#publishing-services-service-types">https://kubernetes.io/docs/concepts/services-networking/service/#publishing-services-service-types</a>	ClusterIP
console.service.port	defines the port number where the web application will be available	80
console.service.serverName	defines the port name where the web application will be available	http
console.service.managementPort	defines the management port used to access metrics and app status	9100
console.resources.requests.memory		1Gi
console.resources.limits.cpu		2
console.resources.limits.memory		2Gi
console.resources.requests.cpu		1
console.env	can be used to set extra environment variables in each app container	[]
console.command	can be used to override the default command provided by the container image. See <a href="https://kubernetes.io/docs/tasks/inject-data-application/define-command-argument-container/">https://kubernetes.io/docs/tasks/inject-data-application/define-command-argument-container/</a>	[]
console.extraVolumes	can be used to define extra volumes for the Console pods, useful for TLS and self-signed certificates	[]
console.extraVolumeMounts	can be used to mount extra volumes for the Console pods, useful for TLS and self-signed certificates	[]
console.startupProbe.enabled	if true, the startup probe is enabled in app container	false
console.startupProbe.scheme	defines the startup probe scheme used on calling the probePath	HTTP
console.startupProbe.probePath	defines the startup probe route used on the app	/health/readiness
console.startupProbe.initialDelaySeconds	defines the number of seconds after the container has started before	30
console.startupProbe.periodSeconds	defines how often the probe is executed	30
console.startupProbe.successThreshold	defines how often it needs to be true to be marked as ready, after failure	1
console.startupProbe.failureThreshold	defines when the probe is considered as failed so the Pod will be marked Unready	5
console.startupProbe.timeoutSeconds	defines the seconds after the probe times out	1
console.readinessProbe.enabled	if true, the readiness probe is enabled in app container	true
console.readinessProbe.scheme	defines the startup probe scheme used on calling the probePath	HTTP
console.readinessProbe.probePath	defines the readiness probe route used on the app	/health/readiness
console.readinessProbe.initialDelaySeconds	defines the number of seconds after the container has started before	30
console.readinessProbe.periodSeconds	defines how often the probe is executed	30
console.readinessProbe.successThreshold	defines how often it needs to be true to be marked as ready, after failure	1
console.readinessProbe.failureThreshold	defines when the probe is considered as failed so the Pod will be marked Unready	5
console.readinessProbe.timeoutSeconds	defines the seconds after the probe times out	1
console.livenessProbe.enabled	if true, the liveness probe is enabled in app container	false
console.livenessProbe.scheme	defines the startup probe scheme used on calling the probePath	HTTP
console.livenessProbe.probePath	defines the liveness probe route used on the app	/health/liveness
console.livenessProbe.initialDelaySeconds	defines the number of seconds after the container has started before	30
console.livenessProbe.periodSeconds	defines how often the probe is executed	30
console.livenessProbe.successThreshold	defines how often it needs to be true to be considered successful after having failed	1
console.livenessProbe.failureThreshold	defines when the probe is considered as failed so the container will be restarted	5
console.livenessProbe.timeoutSeconds	defines the seconds after the probe times out	1
console.metrics.prometheus	Prometheus metrics endpoint	/prometheus
console.serviceAccount.enabled	if true, enables the Console service account	true
console.serviceAccount.name	can be used to set the name of the Console service account	""
console.serviceAccount.annotations	can be used to set the annotations of the Operate service account	{}



Name	Description	Value
console.serviceAccount.automountServiceAccountToken	can be used to control whether the service account token should be automatically mounted	false
console.ingress.enabled	if true, an ingress resource is deployed with the Console deployment. Only useful if an ingress controller is available, like nginx.	false
console.ingress.className	defines the class or configuration of ingress which should be used by the controller	nginx
console.ingress.annotations	defines the ingress related annotations, consumed mostly by the ingress controller	{}
console.ingress.path	defines the path which is associated with the Console service and port <a href="https://kubernetes.io/docs/concepts/services-networking/ingress/#ingress-rules">https://kubernetes.io/docs/concepts/services-networking/ingress/#ingress-rules</a>	/
console.ingress.pathType	can be used to define the Ingress path type. <a href="https://kubernetes.io/docs/concepts/services-networking/ingress/#path-types">https://kubernetes.io/docs/concepts/services-networking/ingress/#path-types</a>	Prefix
console.ingress.host	can be used to define the host of the ingress rule. <a href="https://kubernetes.io/docs/concepts/services-networking/ingress/#ingress-rules">https://kubernetes.io/docs/concepts/services-networking/ingress/#ingress-rules</a>	""
console.ingress.tls.enabled	if true, then tls is configured on the ingress resource. If enabled the Ingress.host need to be defined.	false
console.ingress.tls.secretName	defines the secret name which contains the TLS private key and certificate	camunda-platform-console
console.podSecurityContext	defines the security options the Console broker pod should be run with	
console.podSecurityContext.runAsNonRoot	run as non root	true
console.podSecurityContext.fsGroup		1001
console.containerSecurityContext.allowPrivilegeEscalation		false
console.containerSecurityContext.privileged		false
console.containerSecurityContext.readOnlyRootFilesystem		true
console.containerSecurityContext.runAsNonRoot		true
console.containerSecurityContext.runAsUser		1001
console.nodeSelector	can be used to define on which nodes the Console pods should run	{}
console.tolerations	can be used to define pod toleration's <a href="https://kubernetes.io/docs/concepts/scheduling-eviction/taint-and-toleration/">https://kubernetes.io/docs/concepts/scheduling-eviction/taint-and-toleration/</a>	[]
console.affinity	can be used to define pod affinity or anti-affinity <a href="https://kubernetes.io/docs/concepts/scheduling-eviction/assign-pod-node/#affinity-and-anti-affinity">https://kubernetes.io/docs/concepts/scheduling-eviction/assign-pod-node/#affinity-and-anti-affinity</a>	{}

Zeebe Parameters

Name	Description	Value
zeebe	configuration for the Zeebe sub chart. Contains configuration for the Zeebe broker and related resources.	
zeebe.enabled	if true, all zeebe related resources are deployed via the helm release	true
zeebe.debug	if true, extra info is printed.	false
zeebe.image	configuration to configure the zeebe image specifics	
zeebe.image.registry	can be used to set container image registry.	""
zeebe.image.repository	defines which image repository to use	camunda/zeebe
zeebe.image.tag	can be set to overwrite the global tag, which should be used in that chart	nil
zeebe.image.pullSecrets	can be used to configure image pull secrets <a href="https://kubernetes.io/docs/concepts/containers/images/#specifying-imagepullsecrets-on-a-pod">https://kubernetes.io/docs/concepts/containers/images/#specifying-imagepullsecrets-on-a-pod</a>	[]
zeebe.sidecars	can be used to attach extra containers to the zeebe deployment	[]
zeebe.clusterSize	defines the amount of brokers (=replicas), which are deployed via helm	3
zeebe.partitionCount	defines how many zeebe partitions are set up in the cluster	3
zeebe.replicationFactor	defines how each partition is replicated, the value defines the number of nodes	3
zeebe.env	can be used to set extra environment variables in each zeebe broker container	
zeebe.env[0].name		ZEEBE_BROKER_DATA_SNAPSHOTPERIOD
zeebe.env[0].value		5m
zeebe.env[1].name		ZEEBE_BROKER_DATA_DISK_FREESPACE_REPLICATION
zeebe.env[1].value		2GB
zeebe.env[2].name		ZEEBE_BROKER_DATA_DISK_FREESPACE_PROCESSING
zeebe.env[2].value		3GB
zeebe.configMap	configuration which will be applied to the mounted config map.	
zeebe.configMap.defaultMode	can be used to set permissions on created files by default. Must be an octal value between 0000 and 0777 or a decimal value between 0 and 511. see <a href="https://github.com/kubernetes/api/blob/master/core/v1/types.go#L1615-L1623">https://github.com/kubernetes/api/blob/master/core/v1/types.go#L1615-L1623</a>	754
zeebe.command	can be used to override the default command provided by the container image. See <a href="https://kubernetes.io/docs/tasks/inject-data-application/define-command-argument-container/">https://kubernetes.io/docs/tasks/inject-data-application/define-command-argument-container/</a>	[]

Name	Description	Value
zeebe.logLevel	defines the log level which is used by the zeebe brokers	info
zeebe.log4j2	can be used to overwrite the log4j2 configuration of the zeebe brokers	""
zeebe.javaOpts	can be used to set java options for the zeebe brokers	-XX:+HeapDumpOnOutOfMemoryError -XX:HeapDumpPath=/usr/local/zeebe/data -XX:ErrorFile=/usr/local/zeebe/data/zeebe_error%p.log -XX:+ExitOnOutOfMemoryError
zeebe.service	configuration for the broker service	
zeebe.service.annotations	can be used to define annotations, which will be applied to the Zeebe service	{}
zeebe.service.type	defines the type of the service <a href="https://kubernetes.io/docs/concepts/services-networking/service/#publishing-services-service-types">https://kubernetes.io/docs/concepts/services-networking/service/#publishing-services-service-types</a>	ClusterIP
zeebe.service.httpPort	defines the port of the http endpoint, where for example metrics are provided	9600
zeebe.service.httpName	defines the name of the http endpoint, where for example metrics are provided	http
zeebe.service.commandPort	defines the port of the command api endpoint, where the broker commands are sent to	26501
zeebe.service.commandName	defines the name of the command api endpoint, where the broker commands are sent to	command
zeebe.service.internalPort	defines the port of the internal api endpoint, which is used for internal communication	26502
zeebe.service.internalName	defines the name of the internal api endpoint, which is used for internal communication	internal
zeebe.service.extraPorts	can be used to expose any other ports which are required. Can be useful for exporters	[]
global.zeebe.ServiceAccount	configuration for the service account where the broker pods are assigned to	
zeebe.serviceAccount.enabled	if true, enables the broker service account	true
zeebe.serviceAccount.name	can be used to set the name of the broker service account	""
zeebe.serviceAccount.annotations	can be used to set the annotations of the broker service account	{}
zeebe.serviceAccount.automountServiceAccountToken	can be used to control whether the service account token should be automatically mounted	false
zeebe.cpuThreadCount	defines how many threads can be used for the processing on each broker pod	3
zeebe.ioThreadCount	defines how many threads can be used for the exporting on each broker pod	3
zeebe.resources	configuration to set request and limit configuration for the container <a href="https://kubernetes.io/docs/concepts/configuration/manage-resources-containers/#requests-and-limits">https://kubernetes.io/docs/concepts/configuration/manage-resources-containers/#requests-and-limits</a>	
zeebe.resources.requests		
zeebe.resources.requests.cpu		800m
zeebe.resources.requests.memory		1200Mi
zeebe.resources.limits.cpu		960m
zeebe.resources.limits.memory		1920Mi
zeebe.persistenceType	defines the type of persistence which is used by Zeebe. Possible values are: disk, local and memory.	disk
zeebe.pvcSize	defines the persistent volume claim size, which is used by each broker pod <a href="https://kubernetes.io/docs/concepts/storage/persistent-volumes/#persistentvolumeclaims">https://kubernetes.io/docs/concepts/storage/persistent-volumes/#persistentvolumeclaims</a>	32Gi
zeebe.pvcAccessModes	can be used to configure the persistent volume claim access mode <a href="https://kubernetes.io/docs/concepts/storage/persistent-volumes/#access-modes">https://kubernetes.io/docs/concepts/storage/persistent-volumes/#access-modes</a>	["ReadWriteOnce"]
zeebe.pvcStorageClassName	can be used to set the storage class name which should be used by the persistent volume claim. It is recommended to use a storage class, which is backed with a SSD.	""
zeebe.pvcAnnotations	can be used to specify custom annotations for Zeebe's persistent volume claims, enhancing storage configuration flexibility.	{}
zeebe.extraVolumes	can be used to define extra volumes for the broker pods, useful for additional exporters	[]
zeebe.extraVolumeMounts	can be used to mount extra volumes for the broker pods, useful for additional exporters	[]
zeebe.extraInitContainers	(Deprecated - use <code>initContainers</code> instead) ExtraInitContainers can be used to set up extra init containers for the broker pods, useful for additional exporters	[]
zeebe.initContainers	can be used to set up extra init containers for the application Pod	[]
zeebe.podAnnotations	can be used to define extra broker pod annotations	{}
zeebe.podLabels	can be used to define extra broker pod labels	{}
zeebe.podDisruptionBudget	configuration to configure a pod disruption budget for the broker pods <a href="https://kubernetes.io/docs/tasks/run-application/configure-pdb/">https://kubernetes.io/docs/tasks/run-application/configure-pdb/</a>	
zeebe.podDisruptionBudget.enabled	if true a pod disruption budget is defined for the brokers	false
zeebe.podDisruptionBudget.minAvailable	can be used to set how many pods should be available. Be aware that if minAvailable is set, maxUnavailable will not be set (they are mutually exclusive).	nil
zeebe.podDisruptionBudget.maxUnavailable	can be used to set how many pods should be at max. unavailable	1
zeebe.podSecurityContext	defines the security options the Zeebe broker pod should be run with	

Name	Description	Value
zeebe.podSecurityContext.runAsNonRoot	run as non root	true
zeebe.podSecurityContext.fsGroup		1001
zeebe.containerSecurityContext	defines the security options the Zeebe broker container should be run with	
zeebe.containerSecurityContext.allowPrivilegeEscalation		false
zeebe.containerSecurityContext.privileged		false
zeebe.containerSecurityContext.readOnlyRootFilesystem		true
zeebe.containerSecurityContext.runAsNonRoot		true
zeebe.containerSecurityContext.runAsUser		1001
zeebe.startupProbe	configuration	
zeebe.startupProbe.enabled	if true, the startup probe is enabled in app container	false
zeebe.startupProbe.scheme	defines the startup probe schema used on calling the probePath	HTTP
zeebe.startupProbe.probePath	defines the startup probe route used on the app	/actuator/health/startup
zeebe.startupProbe.initialDelaySeconds	defines the number of seconds after the container has started before the probe is initiated.	30
zeebe.startupProbe.periodSeconds	defines how often the probe is executed	30
zeebe.startupProbe.successThreshold	defines how often it needs to be true to be marked as ready, after failure	1
zeebe.startupProbe.failureThreshold	defines when the probe is considered as failed so the Pod will be marked Unready	5
zeebe.startupProbe.timeoutSeconds	defines the seconds after the probe times out	1
zeebe.readinessProbe	configuration	
zeebe.readinessProbe.enabled	if true, the readiness probe is enabled in app container	true
zeebe.readinessProbe.scheme	defines the startup probe schema used on calling the probePath	HTTP
zeebe.readinessProbe.probePath	defines the readiness probe route used on the app	/actuator/health/readiness
zeebe.readinessProbe.initialDelaySeconds	defines the number of seconds after the container has started before	30
zeebe.readinessProbe.periodSeconds	defines how often the probe is executed	30
zeebe.readinessProbe.successThreshold	defines how often it needs to be true to be marked as ready, after failure	1
zeebe.readinessProbe.failureThreshold	defines when the probe is considered as failed so the Pod will be marked Unready	5
zeebe.readinessProbe.timeoutSeconds	defines the seconds after the probe times out	1
zeebe.livenessProbe	configuration	
zeebe.livenessProbe.enabled	if true, the liveness probe is enabled in app container	false
zeebe.livenessProbe.scheme	defines the startup probe schema used on calling the probePath	HTTP
zeebe.livenessProbe.probePath	defines the liveness probe route used on the app	/actuator/health/liveness
zeebe.livenessProbe.initialDelaySeconds	defines the number of seconds after the container has started before	30
zeebe.livenessProbe.periodSeconds	defines how often the probe is executed	30
zeebe.livenessProbe.successThreshold	defines how often it needs to be true to be considered successful after having failed	1
zeebe.livenessProbe.failureThreshold	defines when the probe is considered as failed so the container will be restarted	5
zeebe.livenessProbe.timeoutSeconds	defines the seconds after the probe times out	1
zeebe.metrics.prometheus	Prometheus metrics endpoint	/actuator/prometheus
zeebe.nodeSelector	can be used to define on which nodes the broker pods should run	{}
zeebe.tolerations	can be used to define pod toleration's	[]
	<a href="https://kubernetes.io/docs/concepts/scheduling-eviction/taint-and-toleration/">https://kubernetes.io/docs/concepts/scheduling-eviction/taint-and-toleration/</a>	
global.zeebe.Affinity	can be used to define pod affinity or anti-affinity	
	<a href="https://kubernetes.io/docs/concepts/scheduling-eviction/assign-pod-node/#affinity-and-anti-affinity">https://kubernetes.io/docs/concepts/scheduling-eviction/assign-pod-node/#affinity-and-anti-affinity</a>	
zeebe.priorityClassName	can be used to define the broker pods priority	""
	<a href="https://kubernetes.io/docs/concepts/scheduling-eviction/pod-priority-preemption/#priorityclass">https://kubernetes.io/docs/concepts/scheduling-eviction/pod-priority-preemption/#priorityclass</a>	
zeebe.retention.enabled	if true, the ILM Policy is created and applied to the index templates.	false
zeebe.retention.minimumAge	defines how old the data must be, before the data is deleted as a duration.	30d
zeebe.retention.policyName	defines the name of the created and applied ILM policy.	zeebe-record-retention-policy
zeebe.configuration	if specified, contents will be used as the application.yaml	""
zeebe.extraConfiguration	if specified, contents will be used for any extra configuration files such as log4j2.xml	{}

ZeebeGateway Parameters

Name	Description	Value
Gateway	configuration to define properties related to the standalone gateway	
zeebeGateway.replicas	defines how many standalone gateways are deployed	2
zeebeGateway.image	configuration to configure the ZeebeGateway image specifics	
zeebeGateway.image.registry	can be used to set container image registry.	""
zeebeGateway.image.repository	defines which image repository to use	camunda/zeebe
zeebeGateway.image.tag	can be set to overwrite the global tag, which should be used in that chart	nil

Name	Description	Value
zeebeGateway.image.pullSecrets	can be used to configure image pull secrets <a href="https://kubernetes.io/docs/concepts/containers/images/#specifying-imagepullsecrets-on-a-pod">https://kubernetes.io/docs/concepts/containers/images/#specifying-imagepullsecrets-on-a-pod</a>	[]
zeebeGateway.sidecars	can be used to attach extra containers to the ZeebeGateway deployment	[]
zeebeGateway.podAnnotations	can be used to define extra gateway pod annotations	{}
zeebeGateway.podLabels	can be used to define extra gateway pod labels	{}
zeebeGateway.logLevel	defines the log level which is used by the gateway	info
zeebeGateway.log4j2	can be used to overwrite the log4j2 configuration of the gateway	""
zeebeGateway.javaOpts	can be used to set java options for the ZeebeGateway	-XX:+ExitOnOutOfMemoryError
zeebeGateway.env	can be used to set extra environment variables in each gateway container	[]
zeebeGateway.configMap	configuration which will be applied to the mounted config map.	
zeebeGateway.configMap.defaultMode	can be used to set permissions on created files by default. Must be an octal value between 0000 and 0777 or a decimal value between 0 and 511.	744
zeebeGateway.command	can be used to override the default command provided by the container image. See <a href="https://kubernetes.io/docs/tasks/inject-data-application/define-command-argument-container/">https://kubernetes.io/docs/tasks/inject-data-application/define-command-argument-container/</a>	[]
zeebeGateway.podDisruptionBudget	configuration to configure a pod disruption budget for the gateway pods <a href="https://kubernetes.io/docs/tasks/run-application/configure-pdb/">https://kubernetes.io/docs/tasks/run-application/configure-pdb/</a>	
zeebeGateway.podDisruptionBudget.enabled	if true a pod disruption budget is defined for the gateways	false
zeebeGateway.podDisruptionBudget.minAvailable	can be used to set how many pods should be available. Be aware that if minAvailable is set, maxUnavailable will not be set (they are mutually exclusive).	1
zeebeGateway.podDisruptionBudget.maxUnavailable	can be used to set how many pods should be at max. unavailable	nil
zeebeGateway.resources	configuration to set request and limit configuration for the container <a href="https://kubernetes.io/docs/concepts/configuration/manage-resources-containers/#requests-and-limits">https://kubernetes.io/docs/concepts/configuration/manage-resources-containers/#requests-and-limits</a>	
zeebeGateway.resources.requests.cpu		400m
zeebeGateway.resources.requests.memory		450Mi
zeebeGateway.resources.limits.cpu		400m
zeebeGateway.resources.limits.memory		450Mi
zeebeGateway.priorityClassName	can be used to define the gateway pods priority <a href="https://kubernetes.io/docs/concepts/scheduling-eviction/pod-priority-preemption/#priorityclass">https://kubernetes.io/docs/concepts/scheduling-eviction/pod-priority-preemption/#priorityclass</a>	""
zeebeGateway.podSecurityContext	defines the security options the gateway pod should be run wit	
zeebeGateway.podSecurityContext.runAsNonRoot		true
zeebeGateway.podSecurityContext.fsGroup		1001
zeebeGateway.containerSecurityContext	defines the security options the gateway container should be run with	
zeebeGateway.containerSecurityContext.allowPrivilegeEscalation		false
zeebeGateway.containerSecurityContext.privileged		false
zeebeGateway.containerSecurityContext.readOnlyRootFilesystem		true
zeebeGateway.containerSecurityContext.runAsNonRoot		true
zeebeGateway.containerSecurityContext.runAsUser		1001
zeebeGateway.startupProbe	configuration	
zeebeGateway.startupProbe.enabled	if true, the startup probe is enabled in app container	false
zeebeGateway.startupProbe.scheme	defines the startup probe schema used on calling the probePath	HTTP
zeebeGateway.startupProbe.probePath	defines the startup probe route used on the app	/actuator/health/startup
zeebeGateway.startupProbe.initialDelaySeconds	defines the number of seconds after the container has started before	30
zeebeGateway.startupProbe.periodSeconds	defines how often the probe is executed	30
zeebeGateway.startupProbe.successThreshold	defines how often it needs to be true to be marked as ready, after failure	1
zeebeGateway.startupProbe.failureThreshold	defines when the probe is considered as failed so the Pod will be marked Unready	5
zeebeGateway.startupProbe.timeoutSeconds	defines the seconds after the probe times out	1
zeebeGateway.readinessProbe	configuration	
zeebeGateway.readinessProbe.enabled	if true, the readiness probe is enabled in app container	true
zeebeGateway.readinessProbe.scheme	defines the startup probe schema used on calling the probePath	HTTP
zeebeGateway.readinessProbe.probePath	defines the readiness probe route used on the app	/actuator/health/readiness
zeebeGateway.readinessProbe.initialDelaySeconds	defines the number of seconds after the container has started before probe is initiated.	30
zeebeGateway.the		
zeebeGateway.readinessProbe.periodSeconds	defines how often the probe is executed	30
zeebeGateway.readinessProbe.successThreshold	defines how often it needs to be true to be marked as ready, after failure	1
zeebeGateway.readinessProbe.failureThreshold	defines when the probe is considered as failed so the Pod will be marked Unready	5
zeebeGateway.readinessProbe.timeoutSeconds	defines the seconds after the probe times out	1
zeebeGateway.livenessProbe	configuration	
zeebeGateway.livenessProbe.enabled	if true, the liveness probe is enabled in app container	false
zeebeGateway.livenessProbe.scheme	defines the startup probe schema used on calling the probePath	HTTP
zeebeGateway.livenessProbe.probePath	defines the liveness probe route used on the app	/actuator/health/liveness

Name	Description	Value
<code>zeebeGateway.livenessProbe.initialDelaySeconds</code>	defines the number of seconds after the container has started before	30
<code>zeebeGateway.livenessProbe.periodSeconds</code>	defines how often the probe is executed	30
<code>zeebeGateway.livenessProbe.successThreshold</code>	defines how often it needs to be true to be considered successful after having failed	1
<code>zeebeGateway.livenessProbe.failureThreshold</code>	defines when the probe is considered as failed so the container will be restarted	5
<code>zeebeGateway.livenessProbe.timeoutSeconds</code>	defines the seconds after the probe times out	1
<code>zeebeGateway.metrics.prometheus</code>	Prometheus metrics endpoint	<code>/actuator/prometheus</code>
<code>zeebeGateway.nodeSelector</code>	can be used to define on which nodes the gateway pods should run	<code>{}</code>
<code>zeebeGateway.tolerations</code>	can be used to define pod toleration's <a href="https://kubernetes.io/docs/concepts/scheduling-eviction/taint-and-toleration/">https://kubernetes.io/docs/concepts/scheduling-eviction/taint-and-toleration/</a>	<code>[]</code>
<code>zeebeGateway.affinity</code>	can be used to define pod affinity or anti-affinity <a href="https://kubernetes.io/docs/concepts/scheduling-eviction/assign-pod-node/#affinity-and-anti-affinity">https://kubernetes.io/docs/concepts/scheduling-eviction/assign-pod-node/#affinity-and-anti-affinity</a>	
<code>zeebeGateway.extraVolumeMounts</code>	can be used to mount extra volumes for the gateway pods, useful for enabling tls between gateway and broker	<code>[]</code>
<code>zeebeGateway.extraVolumes</code>	can be used to define extra volumes for the gateway pods, useful for enabling tls between gateway and broker	<code>[]</code>
<code>zeebeGateway.extraInitContainers</code>	(Deprecated - use <code>initContainers</code> instead) can be used to set up extra init containers for the gateway pods, useful for adding interceptors	<code>[]</code>
<code>zeebeGateway.initContainers</code>	can be used to set up extra init containers for the application Pod	<code>[]</code>
<code>zeebeGateway.service</code>	configuration for the gateway service	
<code>zeebeGateway.service.annotations</code>	can be used to define annotations, which will be applied to the zeebe-gateway service	<code>{}</code>
<code>zeebeGateway.service.type</code>	defines the type of the service <a href="https://kubernetes.io/docs/concepts/services-networking/service/#publishing-services-service-types">https://kubernetes.io/docs/concepts/services-networking/service/#publishing-services-service-types</a>	<code>ClusterIP</code>
<code>zeebeGateway.service.loadBalancerIP</code>	defines public ip of the load balancer if the type is LoadBalancer	<code>""</code>
<code>zeebeGateway.service.loadBalancerSourceRanges</code>	defines list of allowed source ip address ranges if the type is LoadBalancer	<code>[]</code>
<code>zeebeGateway.service.httpPort</code>	defines the port of the http endpoint, where for example metrics are provided	9600
<code>zeebeGateway.service.httpName</code>	defines the name of the http endpoint, where for example metrics are provided	<code>http</code>
<code>zeebeGateway.service.grpcPort</code>	defines the port of the gateway gRPC endpoint, where client commands (grpc) are sent to	26500
<code>zeebeGateway.service.grpcName</code>	defines the name of the gateway gRPC endpoint, where client commands (grpc) are sent to	<code>gateway</code>
<code>zeebeGateway.service.restPort</code>	defines the REST port of the gateway REST endpoint, where client commands (REST) are sent to	8080
<code>zeebeGateway.service.restName</code>	defines the name of the gateway REST endpoint, where client commands (REST) are sent to	<code>rest</code>
<code>zeebeGateway.service.internalPort</code>	defines the port of the internal api endpoint, which is used for internal communication	26502
<code>zeebeGateway.service.internalName</code>	defines the name of the internal api endpoint, which is used for internal communication	<code>internal</code>
<code>zeebeGateway.serviceAccount</code>	configuration for the service account where the gateway pods are assigned to	
<code>zeebeGateway.serviceAccount.enabled</code>	if true, enables the gateway service account	<code>true</code>
<code>zeebeGateway.serviceAccount.name</code>	can be used to set the name of the gateway service account	<code>""</code>
<code>zeebeGateway.serviceAccount.annotations</code>	can be used to set the annotations of the gateway service account	<code>{}</code>
<code>zeebeGateway.serviceAccount.automountServiceAccountToken</code>	can be used to control whether the service account token should be automatically mounted	<code>false</code>
<code>zeebeGateway.ingress.grpc.enabled</code>	if true, an ingress resource is deployed with the Zeebe gateway deployment. Only useful if an ingress controller is available, like nginx.	<code>false</code>
<code>zeebeGateway.ingress.grpc.className</code>	defines the class or configuration of ingress which should be used by the controller	<code>nginx</code>
<code>zeebeGateway.ingress.grpc.annotations</code>	defines the ingress related annotations, consumed mostly by the ingress controller	<code>{}</code>
<code>zeebeGateway.ingress.grpc.path</code>	defines the path which is associated with the Zeebe gateway's gRPC service and port <a href="https://kubernetes.io/docs/concepts/services-networking/ingress/#ingress-rules">https://kubernetes.io/docs/concepts/services-networking/ingress/#ingress-rules</a>	<code>/</code>
<code>zeebeGateway.ingress.grpc.pathType</code>	can be used to define the Ingress path type. <a href="https://kubernetes.io/docs/concepts/services-networking/ingress/#path-types">https://kubernetes.io/docs/concepts/services-networking/ingress/#path-types</a>	<code>Prefix</code>
<code>zeebeGateway.ingress.grpc.host</code>	can be used to define the host of the ingress rule. <a href="https://kubernetes.io/docs/concepts/services-networking/ingress/#ingress-rules">https://kubernetes.io/docs/concepts/services-networking/ingress/#ingress-rules</a>	<code>""</code>
<code>zeebeGateway.ingress.grpc.tls</code>	configuration for tls on the ingress resource <a href="https://kubernetes.io/docs/concepts/services-networking/ingress/#tls">https://kubernetes.io/docs/concepts/services-networking/ingress/#tls</a>	
<code>zeebeGateway.ingress.grpc.tls.enabled</code>	if true, then tls is configured on the ingress resource. If enabled the Ingress.host need to be defined.	<code>false</code>
<code>zeebeGateway.ingress.grpc.tls.secretName</code>	defines the secret name which contains the TLS private key and certificate	<code>camunda-platform-zeebe-gateway-grpc</code>
<code>zeebeGateway.ingress.rest.enabled</code>	if true, an ingress resource is deployed with the Zeebe gateway deployment. Only useful if an ingress controller is available, like nginx.	<code>false</code>
<code>zeebeGateway.ingress.rest.className</code>	defines the class or configuration of ingress which should be used by the controller	<code>nginx</code>
<code>zeebeGateway.ingress.rest.annotations</code>	defines the ingress related annotations, consumed mostly by the ingress controller	<code>{}</code>

Name	Description	Value
zeebeGateway.ingress.rest.path	defines the path which is associated with the Zeebe gateway’s REST service and port <a href="https://kubernetes.io/docs/concepts/services-networking/ingress/#ingress-rules">https://kubernetes.io/docs/concepts/services-networking/ingress/#ingress-rules</a>	/
zeebeGateway.ingress.rest.pathType	can be used to define the Ingress path type. <a href="https://kubernetes.io/docs/concepts/services-networking/ingress/#path-types">https://kubernetes.io/docs/concepts/services-networking/ingress/#path-types</a>	Prefix
zeebeGateway.ingress.rest.host	can be used to define the host of the ingress rule. <a href="https://kubernetes.io/docs/concepts/services-networking/ingress/#ingress-rules">https://kubernetes.io/docs/concepts/services-networking/ingress/#ingress-rules</a>	""
zeebeGateway.ingress.rest.tls	configuration for tls on the ingress resource <a href="https://kubernetes.io/docs/concepts/services-networking/ingress/#tls">https://kubernetes.io/docs/concepts/services-networking/ingress/#tls</a>	
zeebeGateway.ingress.rest.tls.enabled	if true, then tls is configured on the ingress resource. If enabled the Ingress.host need to be defined.	false
zeebeGateway.ingress.rest.tls.secretName	defines the secret name which contains the TLS private key and certificate	camunda-platform-zeebe-gateway-rest
zeebeGateway.contextPath	can be used to make Zeebe web application works on a custom sub-path. This is mainly used to run Camunda web applications under a single domain.	""
zeebeGateway.configuration	if specified, contents will be used as the application.yaml	""
zeebeGateway.extraConfiguration	if specified, contents will be used for any extra configuration files such as log4j2.xml	{}

Operate Parameters

Name	Description	Value
.operate	configuration for the Operate sub chart.	
operate.enabled	if true, the Operate deployment and its related resources are deployed via a helm release	true
operate.image	configuration to configure the Operate image specifics	
operate.image.registry	can be used to set container image registry.	""
operate.image.repository	defines which image repository to use	camunda/operate
operate.image.tag	can be set to overwrite the global tag, which should be used in that chart	nil
operate.image.pullSecrets	can be used to configure image pull secrets <a href="https://kubernetes.io/docs/concepts/containers/images/#specifying-imagepullsecrets-on-a-pod">https://kubernetes.io/docs/concepts/containers/images/#specifying-imagepullsecrets-on-a-pod</a>	[]
operate.sidecars	can be used to attach extra containers to the operate deployment	[]
operate.initContainers	can be used to set up extra init containers for the application Pod	[]
operate.contextPath	can be used to make Operate web application works on a custom sub-path. This is mainly used to run Camunda web applications under a single domain.	""
operate.podAnnotations	can be used to define extra Operate pod annotations	{}
operate.podLabels	can be used to define extra Operate pod labels	{}
operate.logging	configuration for the Operate logging. This template will be directly included in the Operate configuration YAML file	
operate.logging.level.ROOT		INFO
operate.logging.level.io.camunda.operate		INFO
operate.service	configuration to configure the Operate service.	
operate.service.annotations	can be used to define annotations, which will be applied to the Operate service	{}
operate.service.type	defines the type of the service <a href="https://kubernetes.io/docs/concepts/services-networking/service/#publishing-services-service-types">https://kubernetes.io/docs/concepts/services-networking/service/#publishing-services-service-types</a>	ClusterIP
operate.service.port	defines the port of the service, where the Operate web application will be available	80
operate.resources	configuration to set request and limit configuration for the container <a href="https://kubernetes.io/docs/concepts/configuration/manager-resources-containers/#requests-and-limits">https://kubernetes.io/docs/concepts/configuration/manager-resources-containers/#requests-and-limits</a>	
operate.resources.requests.cpu		600m
operate.resources.requests.memory		400Mi
operate.resources.limits.cpu		2000m
operate.resources.limits.memory		2Gi
operate.env	can be used to set extra environment variables in each Operate container	[]
operate.configMap	configuration which will be applied to the mounted config map.	
operate.configMap.defaultMode	can be used to set permissions on created files by default. Must be an octal value between 0000 and 0777 or a decimal value between 0 and 511.	744
operate.command	can be used to override the default command provided by the container image. See <a href="https://kubernetes.io/docs/tasks/inject-data-application/define-command-argument-container/">https://kubernetes.io/docs/tasks/inject-data-application/define-command-argument-container/</a>	[]
operate.extraVolumes	can be used to define extra volumes for the Operate pods, useful for tls and self-signed certificates	[]
operate.extraVolumeMounts	can be used to mount extra volumes for the Operate pods, useful for tls and self-signed certificates	[]

Name	Description	Value
operate.serviceAccount	configuration for the service account where the Operate pods are assigned to	
operate.serviceAccount.enabled	if true, enables the Operate service account	true
operate.serviceAccount.name	can be used to set the name of the Operate service account	""
operate.serviceAccount.annotations	can be used to set the annotations of the Operate service account	{}
operate.serviceAccount.automountServiceAccountToken	can be used to control whether the service account token should be automatically mounted	false
operate.ingress.enabled	if true, an ingress resource is deployed with the Operate deployment. Only useful if an ingress controller is available, like nginx.	false
operate.ingress.className	defines the class or configuration of ingress which should be used by the controller	nginx
operate.ingress.annotations	defines the ingress related annotations, consumed mostly by the ingress controller	{}
operate.ingress.path	defines the path which is associated with the Operate service and port	/
operate.ingress.pathType	<a href="https://kubernetes.io/docs/concepts/services-networking/ingress/#ingress-rules">https://kubernetes.io/docs/concepts/services-networking/ingress/#ingress-rules</a> can be used to define the Ingress path type.	Prefix
operate.ingress.host	<a href="https://kubernetes.io/docs/concepts/services-networking/ingress/#path-types">https://kubernetes.io/docs/concepts/services-networking/ingress/#path-types</a> can be used to define the host of the ingress rule.	""
Ingress.tls	<a href="https://kubernetes.io/docs/concepts/services-networking/ingress/#ingress-rules">https://kubernetes.io/docs/concepts/services-networking/ingress/#ingress-rules</a> configuration for tls on the ingress resource	
operate.ingress.tls.enabled	<a href="https://kubernetes.io/docs/concepts/services-networking/ingress/#tls">https://kubernetes.io/docs/concepts/services-networking/ingress/#tls</a> if true, then tls is configured on the ingress resource. If enabled the Ingress.host need to be defined.	false
operate.ingress.tls.secretName	defines the secret name which contains the TLS private key and certificate	camunda-platform-operate
operate.podSecurityContext	defines the security options the Operate pod should be run with	
operate.podSecurityContext.runAsNonRoot		true
operate.podSecurityContext.fsGroup		1001
operate.containerSecurityContext	defines the security options the Operate container should be run with	
operate.containerSecurityContext.allowPrivilegeEscalation		false
operate.containerSecurityContext.privileged		false
operate.containerSecurityContext.readOnlyRootFilesystem		true
operate.containerSecurityContext.runAsNonRoot		true
operate.containerSecurityContext.runAsUser		1001
operate.startupProbe	configuration	
operate.startupProbe.enabled	if true, the startup probe is enabled in app container	false
operate.startupProbe.scheme	defines the startup probe schema used on calling the probePath	HTTP
operate.startupProbe.probePath	defines the startup probe route used on the app	/actuator/health/readiness
operate.startupProbe.initialDelaySeconds	defines the number of seconds after the container has started before	30
operate.startupProbe.periodSeconds	defines how often the probe is executed	30
operate.startupProbe.successThreshold	defines how often it needs to be true to be marked as ready, after failure	1
operate.startupProbe.failureThreshold	defines when the probe is considered as failed so the Pod will be marked Unready	5
operate.startupProbe.timeoutSeconds	defines the seconds after the probe times out	1
operate.readinessProbe	configuration	
operate.readinessProbe.enabled	if true, the readiness probe is enabled in app container	true
operate.readinessProbe.scheme	defines the startup probe schema used on calling the probePath	HTTP
operate.readinessProbe.probePath	defines the readiness probe route used on the app	/actuator/health/readiness
operate.readinessProbe.initialDelaySeconds	defines the number of seconds after the container has started before	30
operate.readinessProbe.periodSeconds	defines how often the probe is executed	30
operate.readinessProbe.successThreshold	defines how often it needs to be true to be marked as ready, after failure	1
operate.readinessProbe.failureThreshold	defines when the probe is considered as failed so the Pod will be marked Unready	5
operate.readinessProbe.timeoutSeconds	defines the seconds after the probe times out	1
operate.livenessProbe	configuration	
operate.livenessProbe.enabled	if true, the liveness probe is enabled in app container	false
operate.livenessProbe.scheme	defines the startup probe schema used on calling the probePath	HTTP
operate.livenessProbe.probePath	defines the liveness probe route used on the app	/actuator/health/liveness
operate.livenessProbe.initialDelaySeconds	defines the number of seconds after the container has started before	30
operate.livenessProbe.periodSeconds	defines how often the probe is executed	30
operate.livenessProbe.successThreshold	defines how often it needs to be true to be considered successful after having failed	1
operate.livenessProbe.failureThreshold	defines when the probe is considered as failed so the container will be restarted	5
operate.livenessProbe.timeoutSeconds	defines the seconds after the probe times out	1
operate.metrics.prometheus	Prometheus metrics endpoint	/actuator/prometheus
operate.nodeSelector	can be used to define on which nodes the Operate pods should run	{}
operate.tolerations	can be used to define pod toleration's	[]
	<a href="https://kubernetes.io/docs/concepts/scheduling-eviction/taint-and-toleration/">https://kubernetes.io/docs/concepts/scheduling-eviction/taint-and-toleration/</a>	
operate.affinity	can be used to define pod affinity or anti-affinity	{}
	<a href="https://kubernetes.io/docs/concepts/scheduling-eviction/assign-pod-node/#affinity-and-anti-affinity">https://kubernetes.io/docs/concepts/scheduling-eviction/assign-pod-node/#affinity-and-anti-affinity</a>	

Name	Description	Value
<code>operate.retention.enabled</code>	if true, the ILM Policy is created and applied to the index templates.	<code>false</code>
<code>operate.retention.minimumAge</code>	defines how old the data must be, before the data is deleted as a duration.	<code>30d</code>
<code>operate.configuration</code>	if specified, contents will be used as the <code>application.yaml</code>	<code>""</code>
<code>operate.extraConfiguration</code>	if specified, contents will be used for any extra configuration files such as the <code>log4j2.xml</code>	<code>{}</code>

### Tasklist Parameters

Name	Description	Value
<code>tasklist.enabled</code>	if true, the tasklist deployment and its related resources are deployed via a helm release	<code>true</code>
<code>tasklist.image</code>	configuration to configure the tasklist image specifics	
<code>tasklist.image.registry</code>	can be used to set container image registry.	<code>""</code>
<code>tasklist.image.repository</code>	defines which image repository to use	<code>camunda/tasklist</code>
<code>tasklist.image.tag</code>	can be set to overwrite the global tag, which should be used in that chart	<code>nil</code>
<code>tasklist.image.pullSecrets</code>	can be used to configure image pull secrets <a href="https://kubernetes.io/docs/concepts/containers/images/#specifying-imagepullsecrets-on-a-pod">https://kubernetes.io/docs/concepts/containers/images/#specifying-imagepullsecrets-on-a-pod</a>	<code>[]</code>
<code>tasklist.sidecars</code>	can be used to attach extra containers to the tasklist deployment	<code>[]</code>
<code>tasklist.initContainers</code>	can be used to set up extra init containers for the application Pod	<code>[]</code>
<code>tasklist.contextPath</code>	can be used to make Tasklist web application works on a custom sub-path. This is mainly used to run Camunda web applications under a single domain.	<code>""</code>
<code>tasklist.env</code>	can be used to set extra environment variables on each Tasklist container	<code>[]</code>
<code>tasklist.podAnnotations</code>	can be used to define extra Tasklist pod annotations	<code>{}</code>
<code>tasklist.podLabels</code>	can be used to define extra tasklist pod labels	<code>{}</code>
<code>tasklist.configMap</code>	configuration which will be applied to the mounted config map.	
<code>tasklist.configMap.defaultMode</code>	can be used to set permissions on created files by default. Must be an octal value between 0000 and 0777 or a decimal value between 0 and 511.	<code>744</code>
<code>tasklist.command</code>	can be used to override the default command provided by the container image. See <a href="https://kubernetes.io/docs/tasks/inject-data-application/define-command-argument-container/">https://kubernetes.io/docs/tasks/inject-data-application/define-command-argument-container/</a>	<code>[]</code>
<code>tasklist.service</code>	configuration to configure the tasklist service.	
<code>tasklist.service.annotations</code>	can be used to define annotations, which will be applied to the Tasklist service	<code>{}</code>
<code>tasklist.service.type</code>	defines the type of the service <a href="https://kubernetes.io/docs/concepts/services-networking/service/#publishing-services-service-types">https://kubernetes.io/docs/concepts/services-networking/service/#publishing-services-service-types</a>	<code>ClusterIP</code>
<code>tasklist.service.port</code>	defines the port of the service, where the tasklist web application will be available	<code>80</code>
<code>tasklist.identity</code>	configures app user management.	
<code>tasklist.identity.userAccessRestrictions.enabled</code>	if true, enables the identity user access restrictions	<code>true</code>
<code>tasklist.extraVolumes</code>	can be used to define extra volumes for the Tasklist pods, useful for tls and self-signed certificates	<code>[]</code>
<code>tasklist.extraVolumeMounts</code>	can be used to mount extra volumes for the Tasklist pods, useful for tls and self-signed certificates	<code>[]</code>
<code>tasklist.serviceAccount</code>	configuration for the service account where the Tasklist pods are assigned to	
<code>tasklist.serviceAccount.enabled</code>	if true, enables the Tasklist service account	<code>true</code>
<code>tasklist.serviceAccount.name</code>	can be used to set the name of the Tasklist service account	<code>""</code>
<code>tasklist.serviceAccount.annotations</code>	can be used to set the annotations of the Tasklist service account	<code>{}</code>
<code>tasklist.serviceAccount.automountServiceAccountToken</code>	can be used to control whether the service account token should be automatically mounted	<code>false</code>
<code>tasklist.podSecurityContext</code>	defines the security options the Tasklist pod should be run with	
<code>tasklist.podSecurityContext.runAsNonRoot</code>		<code>true</code>
<code>tasklist.podSecurityContext.fsGroup</code>		<code>1001</code>
<code>tasklist.containerSecurityContext</code>	defines the security options the Tasklist container should be run with	
<code>tasklist.containerSecurityContext.allowPrivilegeEscalation</code>		<code>false</code>
<code>tasklist.containerSecurityContext.privileged</code>		<code>false</code>
<code>tasklist.containerSecurityContext.readOnlyRootFilesystem</code>		<code>true</code>
<code>tasklist.containerSecurityContext.runAsNonRoot</code>		<code>true</code>
<code>tasklist.containerSecurityContext.runAsUser</code>		<code>1001</code>
<code>tasklist.startupProbe</code>	configuration	
<code>tasklist.startupProbe.enabled</code>	if true, the startup probe is enabled in app container	<code>false</code>
<code>tasklist.startupProbe.scheme</code>	defines the startup probe schema used on calling the probePath	<code>HTTP</code>
<code>tasklist.startupProbe.probePath</code>	defines the startup probe route used on the app	<code>/actuator/health/readiness</code>
<code>tasklist.startupProbe.initialDelaySeconds</code>	defines the number of seconds after the container has started before	<code>30</code>



Name	Description	Value
<code>tasklist.startupProbe.periodSeconds</code>	defines how often the probe is executed	30
<code>tasklist.startupProbe.successThreshold</code>	defines how often it needs to be true to be marked as ready, after failure	1
<code>tasklist.startupProbe.failureThreshold</code>	defines when the probe is considered as failed so the Pod will be marked Unready	5
<code>tasklist.startupProbe.timeoutSeconds</code>	defines the seconds after the probe times out	1
<code>tasklist.readinessProbe</code>	configuration	
<code>tasklist.readinessProbe.enabled</code>	if true, the readiness probe is enabled in app container	<b>true</b>
<code>tasklist.readinessProbe.scheme</code>	defines the startup probe schema used on calling the probePath	HTTP
<code>tasklist.readinessProbe.probePath</code>	defines the readiness probe route used on the app	<code>/actuator/health/readiness</code>
<code>tasklist.readinessProbe.initialDelaySeconds</code>	defines the number of seconds after the container has started before	30
<code>tasklist.readinessProbe.periodSeconds</code>	defines how often the probe is executed	30
<code>tasklist.readinessProbe.successThreshold</code>	defines how often it needs to be true to be marked as ready, after failure	1
<code>tasklist.readinessProbe.failureThreshold</code>	defines when the probe is considered as failed so the Pod will be marked Unready	5
<code>tasklist.readinessProbe.timeoutSeconds</code>	defines the seconds after the probe times out	1
<code>tasklist.livenessProbe</code>	configuration	
<code>tasklist.livenessProbe.enabled</code>	if true, the liveness probe is enabled in app container	<b>false</b>
<code>tasklist.livenessProbe.scheme</code>	defines the startup probe schema used on calling the probePath	HTTP
<code>tasklist.livenessProbe.probePath</code>	defines the liveness probe route used on the app	<code>/actuator/health/liveness</code>
<code>tasklist.livenessProbe.initialDelaySeconds</code>	defines the number of seconds after the container has started before	30
<code>tasklist.livenessProbe.periodSeconds</code>	defines how often the probe is executed	30
<code>tasklist.livenessProbe.successThreshold</code>	defines how often it needs to be true to be considered successful after having failed	1
<code>tasklist.livenessProbe.failureThreshold</code>	defines when the probe is considered as failed so the container will be restarted	5
<code>tasklist.livenessProbe.timeoutSeconds</code>	defines the seconds after the probe times out	1
<code>tasklist.metrics.prometheus</code>	Prometheus metrics endpoint	<code>/actuator/prometheus</code>
<code>tasklist.nodeSelector</code>	can be used to define on which nodes the Tasklist pods should run	<code>{}</code>
<code>tasklist.tolerations</code>	can be used to define pod toleration's <a href="https://kubernetes.io/docs/concepts/scheduling-eviction/taint-and-toleration/">https://kubernetes.io/docs/concepts/scheduling-eviction/taint-and-toleration/</a>	<code>[]</code>
<code>tasklist.affinity</code>	can be used to define pod affinity or anti-affinity <a href="https://kubernetes.io/docs/concepts/scheduling-eviction/assign-pod-node/#affinity-and-anti-affinity">https://kubernetes.io/docs/concepts/scheduling-eviction/assign-pod-node/#affinity-and-anti-affinity</a>	<code>{}</code>
<code>tasklist.resources</code>	configuration to set request and limit configuration for the container <a href="https://kubernetes.io/docs/concepts/configuration/manage-resources-containers/#requests-and-limits">https://kubernetes.io/docs/concepts/configuration/manage-resources-containers/#requests-and-limits</a>	
<code>tasklist.resources.requests.cpu</code>		400m
<code>tasklist.resources.requests.memory</code>		1Gi
<code>tasklist.resources.limits.cpu</code>		1000m
<code>tasklist.resources.limits.memory</code>		2Gi
<code>tasklist.ingress.enabled</code>	if true, an ingress resource is deployed with the tasklist deployment. Only useful if an ingress controller is available, like nginx.	<b>false</b>
<code>tasklist.ingress.className</code>	defines the class or configuration of ingress which should be used by the controller	<b>nginx</b>
<code>tasklist.ingress.annotations</code>	defines the ingress related annotations, consumed mostly by the ingress controller	<code>{}</code>
<code>tasklist.ingress.path</code>	defines the path which is associated with the operate service and port <a href="https://kubernetes.io/docs/concepts/services-networking/ingress/#ingress-rules">https://kubernetes.io/docs/concepts/services-networking/ingress/#ingress-rules</a>	<code>/</code>
<code>tasklist.ingress.pathType</code>	can be used to define the Ingress path type. <a href="https://kubernetes.io/docs/concepts/services-networking/ingress/#path-types">https://kubernetes.io/docs/concepts/services-networking/ingress/#path-types</a>	<b>Prefix</b>
<code>tasklist.ingress.host</code>	can be used to define the host of the ingress rule. <a href="https://kubernetes.io/docs/concepts/services-networking/ingress/#ingress-rules">https://kubernetes.io/docs/concepts/services-networking/ingress/#ingress-rules</a>	<code>""</code>
<code>tasklist.ingress.tls.enabled</code>	if true, then tls is configured on the ingress resource. If enabled the Ingress.host need to be defined.	<b>false</b>
<code>tasklist.ingress.tls.secretName</code>	defines the secret name which contains the TLS private key and certificate	<b>camunda-platform-tasklist</b>
<code>tasklist.retention.enabled</code>	if true, the ILM Policy is created and applied to the index templates.	<b>false</b>
<code>tasklist.retention.minimumAge</code>	defines how old the data must be, before the data is deleted as a duration.	30d
<code>tasklist.configuration</code>	if specified, contents will be used as the application.yaml	<code>""</code>
<code>tasklist.extraConfiguration</code>	if specified, contents will be used for any extra configuration files such as log4j2.xml	<code>{}</code>

### Optimize Parameters

Name	Description	Value
<code>optimize.enabled</code>	if true, the Optimize deployment and its related resources are deployed via a helm release	<b>true</b>
<code>optimize.image</code>	configuration to configure the Optimize image specifics	
<code>optimize.image.registry</code>	can be used to set container image registry	<code>""</code>
<code>optimize.image.repository</code>	defines which image repository to use	<b>camunda/optimize</b>

Name	Description	Value
optimize.image.tag	can be set to overwrite the global tag, which should be used in that chart	8.5.0
optimize.image.pullSecrets	can be used to configure image pull secrets <a href="https://kubernetes.io/docs/concepts/containers/images/#specifying-imagepullsecrets-on-a-pod">https://kubernetes.io/docs/concepts/containers/images/#specifying-imagepullsecrets-on-a-pod</a>	[]
optimize.migration	configuration for Optimize migration	
optimize.migration.enabled	if true, run Optimize migration script as an init container	true
optimize.migration.env	can be used to set environment variables for Optimize migration init container	[]
optimize.migration.resources	configuration to set request and limit configuration for the migration container <a href="https://kubernetes.io/docs/concepts/configuration/manage-resources-containers/#requests-and-limits">https://kubernetes.io/docs/concepts/configuration/manage-resources-containers/#requests-and-limits</a>	
optimize.migration.resources.requests.cpu		600m
optimize.migration.resources.requests.memory		1Gi
optimize.migration.resources.limits.cpu		2000m
optimize.migration.resources.limits.memory		2Gi
optimize.sidecars	can be used to attach extra containers to the optimize deployment	[]
optimize.contextPath	can be used to make Optimize web application works on a custom sub-path. This is mainly used to run Camunda web applications under a single domain.	""
optimize.configMap	configuration which will be applied to the mounted config map.	
optimize.configMap.defaultMode	can be used to set permissions on created files by default. Must be an octal value between 0000 and 0777 or a decimal value between 0 and 511.	754
optimize.podAnnotations	can be used to define extra Optimize pod annotations	{}
optimize.podLabels	can be used to define extra Optimize pod labels	{}
optimize.partitionCount	defines how many Zeebe partitions are set up in the cluster and which should be imported by Optimize	3
optimize.env	can be used to set extra environment variables in each Optimize container	[]
optimize.command	can be used to override the default command provided by the container image. See <a href="https://kubernetes.io/docs/tasks/inject-data-application/define-command-argument-container/">https://kubernetes.io/docs/tasks/inject-data-application/define-command-argument-container/</a>	[]
optimize.extraVolumes	can be used to define extra volumes for the Optimize pods, useful for tls and self-signed certificates	[]
optimize.extraVolumeMounts	can be used to mount extra volumes for the Optimize pods, useful for tls and self-signed certificates	[]
optimize.initContainers	can be used to set up extra init containers for the application Pod	[]
optimize.serviceAccount	configuration for the service account where the Optimize pods are assigned to	
optimize.serviceAccount.enabled	if true, enables the Optimize service account	true
optimize.serviceAccount.name	can be used to set the name of the Optimize service account	""
optimize.serviceAccount.annotations	can be used to set the annotations of the Optimize service account	{}
optimize.serviceAccount.automountServiceAccountToken	can be used to control whether the service account token should be automatically mounted	false
optimize.service	configuration to configure the Optimize service.	
optimize.service.annotations	can be used to define annotations, which will be applied to the Optimize service	{}
optimize.service.type	defines the type of the service <a href="https://kubernetes.io/docs/concepts/services-networking/service/#publishing-services-service-types">https://kubernetes.io/docs/concepts/services-networking/service/#publishing-services-service-types</a>	ClusterIP
optimize.service.port	defines the port of the service, where the Optimize web application will be available	80
optimize.service.managementPort	defines the port where actuator will be available. Also required to reach backup API	8092
optimize.podSecurityContext	defines the security options the Optimize pod should be run with	
optimize.podSecurityContext.runAsNonRoot		true
optimize.podSecurityContext.fsGroup		1001
optimize.containerSecurityContext	defines the security options the Optimize container should be run with	
optimize.containerSecurityContext.allowPrivilegeEscalation		false
optimize.containerSecurityContext.privileged		false
optimize.containerSecurityContext.readOnlyRootFilesystem		true
optimize.containerSecurityContext.runAsNonRoot		true
optimize.containerSecurityContext.runAsUser		1001
optimize.startupProbe	configuration	
optimize.startupProbe.enabled	if true, the startup probe is enabled in app container	false
optimize.startupProbe.scheme	defines the startup probe schema used on calling the probePath	HTTP
optimize.startupProbe.probePath	defines the startup probe route used on the app	/api/readyz
optimize.startupProbe.initialDelaySeconds	defines the number of seconds after the container has started before	30
optimize.startupProbe.periodSeconds	defines how often the probe is executed	30
optimize.startupProbe.successThreshold	defines how often it needs to be true to be marked as ready, after failure	1
optimize.startupProbe.failureThreshold	defines when the probe is considered as failed so the Pod will be marked Unready	5
optimize.startupProbe.timeoutSeconds	defines the seconds after the probe times out	1
optimize.readinessProbe	configuration	

Name	Description	Value
optimize.readinessProbe.enabled	if true, the readiness probe is enabled in app container	true
optimize.readinessProbe.scheme	defines the startup probe schema used on calling the probePath	HTTP
optimize.readinessProbe.probePath	defines the readiness probe route used on the app	/api/readyz
optimize.readinessProbe.initialDelaySeconds	defines the number of seconds after the container has started before	30
optimize.readinessProbe.periodSeconds	defines how often the probe is executed	30
optimize.readinessProbe.successThreshold	defines how often it needs to be true to be marked as ready, after failure	1
optimize.readinessProbe.failureThreshold	defines when the probe is considered as failed so the Pod will be marked Unready	5
optimize.readinessProbe.timeoutSeconds	defines the seconds after the probe times out	1
optimize.livenessProbe	configuration	
optimize.livenessProbe.enabled	if true, the liveness probe is enabled in app container	false
optimize.livenessProbe.scheme	defines the startup probe schema used on calling the probePath	HTTP
optimize.livenessProbe.probePath	defines the liveness probe route used on the app	/api/readyz
optimize.livenessProbe.initialDelaySeconds	defines the number of seconds after the container has started before	30
optimize.livenessProbe.periodSeconds	defines how often the probe is executed	30
optimize.livenessProbe.successThreshold	defines how often it needs to be true to be considered successful after having failed	1
optimize.livenessProbe.failureThreshold	defines when the probe is considered as failed so the container will be restarted	5
optimize.livenessProbe.timeoutSeconds	defines the seconds after the probe times out	1
optimize.metrics.prometheus	Prometheus metrics endpoint	/actuator/prometheus
optimize.nodeSelector	can be used to define on which nodes the Optimize pods should run	{}
optimize.tolerations	can be used to define pod toleration's	[]
optimize.affinity	https://kubernetes.io/docs/concepts/scheduling-eviction/taint-and-toleration/ can be used to define pod affinity or anti-affinity https://kubernetes.io/docs/concepts/scheduling-eviction/assign-pod-node/#affinity-and-anti-affinity	{}
optimize.resources	configuration to set request and limit configuration for the container https://kubernetes.io/docs/concepts/configuration/manage-resources-containers/#requests-and-limits	
optimize.resources.requests.cpu		600m
optimize.resources.requests.memory		1Gi
optimize.resources.limits.cpu		2000m
optimize.resources.limits.memory		2Gi
optimize.ingress.enabled	if true, an ingress resource is deployed with the Optimize deployment. Only useful if an ingress controller is available, like nginx.	false
optimize.ingress.className	defines the class or configuration of ingress which should be used by the controller	nginx
optimize.ingress.annotations	defines the ingress related annotations, consumed mostly by the ingress controller	{}
optimize.ingress.path	defines the path which is associated with the operate service and port https://kubernetes.io/docs/concepts/services-networking/ingress/#ingress-rules	/
optimize.ingress.pathType	can be used to define the Ingress path type. https://kubernetes.io/docs/concepts/services-networking/ingress/#path-types	Prefix
optimize.ingress.host	can be used to define the host of the ingress rule. https://kubernetes.io/docs/concepts/services-networking/ingress/#ingress-rules	""
optimize.ingress.tls	configuration for tls on the ingress resource https://kubernetes.io/docs/concepts/services-networking/ingress/#tls	
optimize.ingress.tls.enabled	if true, then tls is configured on the ingress resource. If enabled the Ingress.host need to be defined.	false
optimize.ingress.tls.secretName	defines the secret name which contains the TLS private key and certificate	camunda-platform-optimize
optimize.configuration	if specified, contents will be used as the environment-config.yaml	""
optimize.extraConfiguration	if specified, contents will be used for any extra configuration files such as environment-logback.xml	{}

### Identity Parameters

Name	Description	Value
identity.enabled	if true, the identity deployment and its related resources are deployed via a helm release	true
identity.fullnameOverride	can be used to override the full name of the Identity resources	""
identity.nameOverride	can be used to partly override the name of the Identity resources (names will still be prefixed with the release name)	""
identity.firstUser	configuration to configure properties of the first Identity user, which can be used to access all	
identity.firstUser.enabled	if true, Identity will seed the first user in Keycloak.	true
identity.firstUser.username	defines the username of the first user, needed to log in into the web applications	demo

Name	Description	Value
identity.firstUser.password	defines the password of the first user, needed to log in into the web applications	demo
identity.firstUser.email	defines the email address of the first user; a valid email address is required to use WebModeler	demo@example.org
identity.firstUser.firstName	defines the first name of the first user; a name is required to use WebModeler	Demo
identity.firstUser.lastName	defines the last name of the first user; a name is required to use WebModeler	User
identity.firstUser.existingSecret	can be used to use an own existing secret for Identity first user.	""
identity.image	configuration to configure the identity image specifics	""
identity.image.registry	can be used to set container image registry.	""
identity.image.repository	defines which image repository to use	camunda/identity
identity.image.tag	can be set to overwrite the global tag, which should be used in that chart	nil
identity.image.pullSecrets	can be used to configure image pull secrets	[]
	<a href="https://kubernetes.io/docs/concepts/containers/images/#specifying-imagepullsecrets-on-a-pod">https://kubernetes.io/docs/concepts/containers/images/#specifying-imagepullsecrets-on-a-pod</a>	
identity.sidecars	can be used to attach extra containers to the identity deployment	[]
identity.initContainers	can be used to set up extra init containers for the application Pod	[]
identity.fullURL	can be used when Ingress is configured (for both multi and single domain setup).	""
identity.contextPath	can be used to make Identity web application works on a custom sub-path. This is mainly used to run Camunda web applications under a single domain.	""
identity.podAnnotations	can be used to define extra Identity pod annotations	{}
identity.podLabels	can be used to define extra Identity pod labels	{}
identity.service	configuration to configure the identity service.	
identity.service.annotations	can be used to define annotations, which will be applied to the identity service	{}
identity.service.type	defines the type of the service <a href="https://kubernetes.io/docs/concepts/services-networking/service/#publishing-services-service-types">https://kubernetes.io/docs/concepts/services-networking/service/#publishing-services-service-types</a>	ClusterIP
identity.service.port	defines the port of the service on which the identity application will be available	80
identity.service.metricsPort	defines the port of the service on which the identity metrics will be available	82
identity.service.metricsName	defines the name of the service on which the identity metrics will be available	metrics
identity.podSecurityContext	defines the security options the Identity pod should be run with	
identity.podSecurityContext.runAsNonRoot		true
identity.podSecurityContext.fsGroup		1001
identity.containerSecurityContext	defines the security options the Identity container should be run with	
identity.containerSecurityContext.allowPrivilegeEscalation		false
identity.containerSecurityContext.privileged		false
identity.containerSecurityContext.readOnlyRootFilesystem		true
identity.containerSecurityContext.runAsNonRoot		true
identity.containerSecurityContext.runAsUser		1001
identity.startupProbe	configuration	
identity.startupProbe.enabled	if true, the startup probe is enabled in app container	false
identity.startupProbe.scheme	defines the startup probe schema used on calling the probePath	HTTP
identity.startupProbe.probePath	defines the startup probe route used on the app	/actuator/health
identity.startupProbe.initialDelaySeconds	defines the number of seconds after the container has started before the probe is initiated.	30
identity.startupProbe.periodSeconds	defines how often the probe is executed	30
identity.startupProbe.successThreshold	defines how often it needs to be true to be marked as ready, after failure	1
identity.startupProbe.failureThreshold	defines when the probe is considered as failed so the Pod will be marked Unready	5
identity.startupProbe.timeoutSeconds	defines the seconds after the probe times out	1
identity.readinessProbe	configuration	
identity.readinessProbe.enabled	if true, the readiness probe is enabled in app container	true
identity.readinessProbe.scheme	defines the startup probe schema used on calling the probePath	HTTP
identity.readinessProbe.probePath	defines the readiness probe route used on the app	/actuator/health
identity.readinessProbe.initialDelaySeconds	defines the number of seconds after the container has started before the probe is initiated.	30
identity.readinessProbe.periodSeconds	defines how often the probe is executed	30
identity.readinessProbe.successThreshold	defines how often it needs to be true to be marked as ready, after failure	1
identity.readinessProbe.failureThreshold	defines when the probe is considered as failed so the Pod will be marked Unready	5
identity.readinessProbe.timeoutSeconds	defines the seconds after the probe times out	1
identity.livenessProbe	configuration	
identity.livenessProbe.enabled	if true, the liveness probe is enabled in app container	false
identity.livenessProbe.scheme	defines the startup probe schema used on calling the probePath	HTTP
identity.livenessProbe.probePath	defines the liveness probe route used on the app	/actuator/health
identity.livenessProbe.initialDelaySeconds	defines the number of seconds after the container has started before	30
identity.livenessProbe.periodSeconds	defines how often the probe is executed	30
identity.livenessProbe.successThreshold	defines how often it needs to be true to be considered successful after having failed	1
identity.livenessProbe.failureThreshold	defines when the probe is considered as failed so the container will be restarted	5

Name	Description	Value
identity.livenessProbe.timeoutSeconds	defines the seconds after the probe times out	1
identity.metrics.prometheus	Prometheus metrics endpoint	/actuator/prometheus
identity.nodeSelector	can be used to define on which nodes the Identity pods should run	{}
identity.tolerations	can be used to define pod toleration's	[]
identity.affinity	https://kubernetes.io/docs/concepts/scheduling-eviction/taint-and-toleration/ can be used to define pod affinity or anti-affinity https://kubernetes.io/docs/concepts/scheduling-eviction/assign-pod-node/#affinity-and-anti-affinity	{}
identity.resources	configuration to set request and limit configuration for the container https://kubernetes.io/docs/concepts/configuration/manage-resources-containers/#requests-and-limits	
identity.resources.requests.memory		400Mi
identity.resources.limits.cpu		2000m
identity.resources.requests.cpu		600m
identity.resources.limits.memory		2Gi
identity.env	can be used to set extra environment variables in each identity container. See the documentation https://docs.camunda.io/docs/self-managed/identity/deployment/configuration-variables/ for more details.	[]
identity.command	can be used to override the default command provided by the container image. See https://kubernetes.io/docs/tasks/inject-data-application/define-command-argument-container/	[]
identity.extraVolumes	can be used to define extra volumes for the identity pods, useful for tls and self-signed certificates	[]
identity.extraVolumeMounts	can be used to mount extra volumes for the identity pods, useful for tls and self-signed certificates	[]
identity.serviceAccount	configuration for the service account where the identity pods are assigned to	
identity.serviceAccount.enabled	if true, enables the identity service account	true
identity.serviceAccount.name	can be used to set the name of the identity service account	""
identity.serviceAccount.annotations	can be used to set the annotations of the identity service account	{}
identity.serviceAccount.automountServiceAccountToken	can be used to control whether the service account token should be automatically mounted	true
identity.ingress.enabled	if true, an ingress resource is deployed with the identity deployment. Only useful if an ingress controller is available, like nginx.	false
identity.ingress.className	defines the class or configuration of ingress which should be used by the controller	nginx
identity.ingress.annotations	defines the ingress related annotations, consumed mostly by the ingress controller	{}
identity.ingress.path	defines the path which is associated with the operate service and port https://kubernetes.io/docs/concepts/services-networking/ingress/#ingress-rules	/
identity.ingress.pathType	can be used to define the Ingress path type. https://kubernetes.io/docs/concepts/services-networking/ingress/#path-types	Prefix
identity.ingress.host	can be used to define the host of the ingress rule. https://kubernetes.io/docs/concepts/services-networking/ingress/#ingress-rules	""
identity.ingress.tls	configuration for tls on the ingress resource https://kubernetes.io/docs/concepts/services-networking/ingress/#tls	
identity.ingress.tls.enabled	if true, then tls is configured on the ingress resource. If enabled the Ingress.host need to be defined.	false
identity.ingress.tls.secretName	defines the secret name which contains the TLS private key and certificate	camunda-platform-identity
identity.externalDatabase.enabled		false
identity.externalDatabase.host	Database host	nil
identity.externalDatabase.port	Database port number	nil
identity.externalDatabase.username	Non-root username	nil
identity.externalDatabase.password	Password for the non-root username	nil
identity.externalDatabase.database	The database name	nil
identity.externalDatabase.existingSecret	Name of an existing secret resource containing the database credentials	nil
identity.externalDatabase.existingSecretPasswordKey	Name of an existing secret key containing the database credentials	nil
identity.configuration	if specified, contents will be used as the application.yaml	""
identity.extraConfiguration	if specified, contents will be used for any extra configuration files such as the log4j2.xml	{}

#### Identity - PostgreSQL Parameters

Name	Description	Value
identityPostgresql	configuration for the PostgreSQL dependency chart used by Identity. For more details, check Bitnami package for PostgreSQL documentation.	
identityPostgresql.enabled	Enable Identity PostgreSQL Helm chart. Required for Multi-Tenancy.	false
identityPostgresql.image.repository	PostgreSQL repo	bitnami/postgresql
identityPostgresql.image.tag	PostgreSQL image tag	15.6.0
identityPostgresql.nameOverride	the name used for Identity PostgreSQL.	identity-postgresql
identityPostgresql.auth.username	Non-root username	identity
identityPostgresql.auth.database	The database name	identity
identityPostgresql.auth.password	Password for the non-root username	nil
identityPostgresql.auth.existingSecret	Name of an existing secret resource containing the database credentials	nil

Identity - Keycloak Parameters

Name	Description	Value
identityKeycloak	configuration, for the Keycloak dependency chart which is used by Identity. For more details, check Bitnami package for Keycloak documentation.	
identityKeycloak.enabled	Enable Identity Keycloak Helm chart. It is used incorporate with “global.identity.keycloak” to use your own Keycloak instead of the one comes with Camunda Helm chart	true
identityKeycloak.nameOverride	the name used for Keycloak.	keycloak
identityKeycloak.image	configuration.	
identityKeycloak.image.repository	image repo	bitnami/keycloak
identityKeycloak.image.tag	image tag	23.0.7
identityKeycloak.postgresql	configuration.	
identityKeycloak.postgresql.image.repository	image repo	bitnami/postgresql
identityKeycloak.postgresql.image.tag	image tag	15.6.0
identityKeycloak.postgresql.primary.containerSecurityContext.enabled		true
identityKeycloak.postgresql.primary.containerSecurityContext.privileged		false
identityKeycloak.postgresql.primary.containerSecurityContext.readOnlyRootFilesystem		true
identityKeycloak.postgresql.primary.containerSecurityContext.allowPrivilegeEscalation		false
identityKeycloak.postgresql.primary.containerSecurityContext.runAsNonRoot		true
identityKeycloak.postgresql.primary.containerSecurityContext.runAsUser		1001
identityKeycloak.postgresql.primary.containerSecurityContext.capabilities.drop		["ALL"]
identityKeycloak.postgresql.primary.containerSecurityContext.seccompProfile.type		RuntimeDefault
identityKeycloak.postgresql.primary.podSecurityContext.enabled		true
identityKeycloak.postgresql.primary.podSecurityContext.runAsNonRoot		true
identityKeycloak.postgresql.primary.podSecurityContext.fsGroup		1001
identityKeycloak.proxy	keycloak proxy	edge
identityKeycloak.tls	can be used to enable TLS encryption. Required for HTTPs traffic.	
identityKeycloak.tls.enabled	enabling tls	false
identityKeycloak.initContainers[0].name		copy-camunda-theme
identityKeycloak.initContainers[0].image		{{ .Values.global.identity.image   default "camunda/identity:latest" }}
identityKeycloak.initContainers[0].imagePullPolicy		{{ .Values.global.identity.imagePullPolicy   default "Always" }}
identityKeycloak.initContainers[0].command		["sh","-c","cp -a /app/keycloak-theme/* /mnt"]
identityKeycloak.initContainers[0].securityContext.privileged		false
identityKeycloak.initContainers[0].securityContext.readOnlyRootFilesystem		true
identityKeycloak.initContainers[0].securityContext.allowPrivilegeEscalation		false
identityKeycloak.initContainers[0].securityContext.runAsNonRoot		true
identityKeycloak.initContainers[0].securityContext.runAsUser		1001
identityKeycloak.initContainers[0].securityContext.capabilities.drop		["ALL"]
identityKeycloak.initContainers[0].securityContext.seccompProfile.type		RuntimeDefault
identityKeycloak.initContainers[0].volumeMounts[0].name		camunda-theme
identityKeycloak.initContainers[0].volumeMounts[0].mountPath		/mnt
identityKeycloak.extraVolumeMounts[0].name		camunda-theme
identityKeycloak.extraVolumeMounts[0].mountPath		/opt/bitnami/keycloak/themes/identity
identityKeycloak.extraVolumeMounts[1].mountPath		/opt/bitnami/keycloak/data/tmp
identityKeycloak.extraVolumeMounts[1].name		data-tmp
identityKeycloak.containerSecurityContext.privileged		false
identityKeycloak.containerSecurityContext.readOnlyRootFilesystem		true
identityKeycloak.containerSecurityContext.allowPrivilegeEscalation		false
identityKeycloak.containerSecurityContext.runAsNonRoot		true

Name	Description	Value
identityKeycloak.containerSecurityContext.runAsUser	defines the context for Keycloak. This config is valid for Keycloak v19.x.x only	1001
identityKeycloak.containerSecurityContext.capabilities.drop		["ALL"]
identityKeycloak.containerSecurityContext.seccompProfile.type		RuntimeDefault
identityKeycloak.podSecurityContext.runAsNonRoot		true
identityKeycloak.podSecurityContext.fsGroup		1001
identityKeycloak.httpRelativePath		/auth/
identityKeycloak.extraEnvVars		
identityKeycloak.extraEnvVars[0].name		KEYCLOAK_PROXY_ADDRESS_FORWARDING
identityKeycloak.extraEnvVars[0].value		{{ .Values.global.ingress.tls.enabled }}
identityKeycloak.ingress.enabled		false
identityKeycloak.ingress.tls	can be used to enable TLS configuration for the host defined at ingress.hostname parameter.	false
identityKeycloak.ingress.extraTls	configuration for additional hostnames to be covered with this ingress record.	[]
identityKeycloak.ingress.annotations	configures annotations to be applied to the ingress record.	{}
identityKeycloak.service	configuration, to configure the service which is deployed along with keycloak	
identityKeycloak.service.type	can be set to change the service type.	ClusterIP
identityKeycloak.auth	uses the secrets generated by keycloak, to access keycloak.	
identityKeycloak.auth.adminUser	defines the keycloak administrator user	admin
identityKeycloak.auth.existingSecret	can be used to reuse an existing secret containing authentication information.	""

WebModeler Parameters

Name	Description	Value
webModeler.enabled	if true, the WebModeler deployment and its related resources are deployed via a helm release	false
webModeler.fullnameOverride	can be used to override the full name of the WebModeler resources	""
webModeler.nameOverride	can be used to partly override the name of the WebModeler resources (names will still be prefixed with the release name)	""
webModeler.image	configuration of the WebModeler Docker images	
webModeler.image.registry	can be used to set the Docker registry for the WebModeler images (overwrites global.image.registry)	registry.camunda.cloud
webModeler.image.tag	can be used to set the Docker image tag for the WebModeler images (overwrites global.image.tag)	8.5.0
webModeler.image.pullSecrets	can be used to configure image pull secrets, see <a href="https://kubernetes.io/docs/concepts/containers/images/#specifying-imagepullsecrets-on-a-pod">https://kubernetes.io/docs/concepts/containers/images/#specifying-imagepullsecrets-on-a-pod</a>	[]
webModeler.contextPath	can be used to make WebModeler available on a custom sub-path. This is mainly used to run the Camunda web applications under a single domain.	""

WebModeler - RestAPI Parameters

Name	Description	Value
webModeler.restapi	configuration of the WebModeler restapi component	
webModeler.restapi.image	configuration of the restapi Docker image	
webModeler.restapi.image.repository	defines which image repository to use for the restapi Docker image	web-modeler-ee/modeler-restapi
webModeler.restapi.sidecars	can be used to attach extra containers to the restapi deployment	[]
webModeler.restapi.initContainers	can be used to set up extra init containers for the application Pod	[]
webModeler.restapi.externalDatabase	can be used to configure a connection to an external database. This will only be applied	
webModeler.restapi.externalDatabase.url	defines the JDBC url of the database instance	""
webModeler.restapi.externalDatabase.user		""
webModeler.restapi.externalDatabase.password	defines the database user's password	""
webModeler.restapi.mail	configuration for emails sent by WebModeler	
webModeler.restapi.mail.smtpHost	defines the host name of the SMTP server to be used by WebModeler	""
webModeler.restapi.mail.smtpPort	defines the port number of the SMTP server	587
webModeler.restapi.mail.smtpUser	can be used to provide a user for the SMTP server	""
webModeler.restapi.mail.smtpPassword	can be used to provide a password for the SMTP server	""
webModeler.restapi.mail.smtpTlsEnabled	if true, enforces TLS encryption for SMTP connections (using STARTTLS)	true
webModeler.restapi.mail.fromAddress	defines the email address that will be displayed as the sender of emails sent by WebModeler	""
webModeler.restapi.mail.fromName	defines the name that will be displayed as the sender of emails sent by WebModeler	Camunda 8

Name	Description	Value
webModeler.restapi.podAnnotations	can be used to define extra restapi pod annotations	{}
webModeler.restapi.podLabels	can be used to define extra restapi pod labels	{}
webModeler.restapi.env	can be used to set extra environment variables in each restapi container	[]
webModeler.restapi.command	can be used to override the default command provided by the container image, see <a href="https://kubernetes.io/docs/tasks/inject-data-application/define-command-argument-container/">https://kubernetes.io/docs/tasks/inject-data-application/define-command-argument-container/</a>	[]
webModeler.restapi.extraVolumes	can be used to define extra volumes for the restapi pods, useful for TLS and self-signed certificates	[]
webModeler.restapi.extraVolumeMounts	can be used to mount extra volumes for the restapi pods, useful for TLS and self-signed certificates	[]
webModeler.restapi.podSecurityContext	can be used to define the security options the restapi pod should be run with	
webModeler.restapi.podSecurityContext.runAsNonRoot		true
webModeler.restapi.podSecurityContext.fsGroup		1001
webModeler.restapi.containerSecurityContext		
webModeler.restapi.containerSecurityContext.privileged		false
webModeler.restapi.containerSecurityContext.readOnlyRootFilesystem	can be used to define the security options the restapi container should be run with	true
webModeler.restapi.containerSecurityContext.allowPrivilegeEscalation		false
webModeler.restapi.containerSecurityContext.runAsNonRoot		true
webModeler.restapi.containerSecurityContext.runAsUser		1001
webModeler.restapi.startupProbe	configuration of the restapi startup probe	
webModeler.restapi.startupProbe.enabled	if true, the startup probe will be enabled for the restapi container	false
webModeler.restapi.startupProbe.scheme	defines the startup probe schema used on calling the probePath	HTTP
webModeler.restapi.startupProbe.probePath	defines the HTTP endpoint used for the startup probe	/health/liveness
webModeler.restapi.startupProbe.initialDelaySeconds	defines the number of seconds after the container has started before the probe is initiated	30
webModeler.restapi.startupProbe.periodSeconds	defines how often the probe is executed	30
webModeler.restapi.startupProbe.successThreshold	defines how often the probe needs to succeed to be considered successful after having failed	1
webModeler.restapi.startupProbe.failureThreshold	defines when the probe is considered failed so the container will be restarted	5
webModeler.restapi.startupProbe.timeoutSeconds	defines the number of seconds after which the probe times out	1
webModeler.restapi.readinessProbe	configuration of the restapi readiness probe	
webModeler.restapi.readinessProbe.enabled	if true, the readiness probe will be enabled for the restapi container	true
webModeler.restapi.readinessProbe.scheme	defines the startup probe schema used on calling the probePath	HTTP
webModeler.restapi.readinessProbe.probePath	defines the HTTP endpoint used for the readiness probe	/health/readiness
webModeler.restapi.readinessProbe.initialDelaySeconds	defines the number of seconds after the container has started before the probe is initiated	30
webModeler.restapi.readinessProbe.periodSeconds	defines how often the probe is executed	30
webModeler.restapi.readinessProbe.successThreshold	defines how often the probe needs to succeed to be considered successful after having failed	1
webModeler.restapi.readinessProbe.failureThreshold	defines when the probe is considered failed so the Pod will be marked unready	5
webModeler.restapi.readinessProbe.timeoutSeconds	defines the number of seconds after which the probe times out	1
webModeler.restapi.livenessProbe	configuration of the restapi liveness probe	
webModeler.restapi.livenessProbe.enabled	if true, the liveness probe will be enabled for the restapi container	false
webModeler.restapi.livenessProbe.scheme	defines the startup probe schema used on calling the probePath	HTTP
webModeler.restapi.livenessProbe.probePath	defines the HTTP endpoint used for the liveness probe	/health/liveness
webModeler.restapi.livenessProbe.initialDelaySeconds	defines the number of seconds after the container has started before the probe is initiated	30
webModeler.restapi.livenessProbe.periodSeconds	defines how often the probe is executed	30
webModeler.restapi.livenessProbe.successThreshold	defines how often the probe needs to succeed to be considered successful after having failed	1
webModeler.restapi.livenessProbe.failureThreshold	defines when the probe is considered failed so the container will be restarted	5
webModeler.restapi.livenessProbe.timeoutSeconds	defines the number of seconds after which the probe times out	1
webModeler.restapi.metrics.prometheus	Prometheus metrics endpoint	/metrics
webModeler.restapi.nodeSelector	can be used to select the nodes the restapi pods should run on	{}
webModeler.restapi.tolerations	can be used to define pod tolerations, see <a href="https://kubernetes.io/docs/concepts/scheduling-eviction/taint-and-toleration/">https://kubernetes.io/docs/concepts/scheduling-eviction/taint-and-toleration/</a>	[]
webModeler.restapi.affinity	can be used to define pod affinity or anti-affinity, see <a href="https://kubernetes.io/docs/concepts/scheduling-eviction/assign-pod-node/#affinity-and-anti-affinity">https://kubernetes.io/docs/concepts/scheduling-eviction/assign-pod-node/#affinity-and-anti-affinity</a>	{}
webModeler.restapi.resources	configuration of resource requests and limits for the container, see <a href="https://kubernetes.io/docs/concepts/configuration/manage-resources-containers/#requests-and-limits">https://kubernetes.io/docs/concepts/configuration/manage-resources-containers/#requests-and-limits</a>	
webModeler.restapi.resources.requests.cpu		500m
webModeler.restapi.resources.requests.memory		1Gi



Name	Description	Value
webModeler.restapi.resources.limits.cpu		1000m
webModeler.restapi.resources.limits.memory		2Gi
webModeler.restapi.service	configuration of the WebModeler restapi service	
webModeler.restapi.service.annotations	can be used to define annotations which will be applied to the service	{}
webModeler.restapi.service.type	defines the type of the service, see <a href="https://kubernetes.io/docs/concepts/services-networking/service/#publishing-services-service-types">https://kubernetes.io/docs/concepts/services-networking/service/#publishing-services-service-types</a>	ClusterIP
webModeler.restapi.service.port	defines the default port of the service	80
webModeler.restapi.service.managementPort	defines the management port of the service	8091
webModeler.restapi.configuration	if specified, contents will be used as the application.yaml	""
webModeler.restapi.extraConfiguration	if specified, contents will be used for any extra configuration files such as log4j2.xml	{}

WebModeler - WebApp Parameters

Name	Description	Value
webModeler.webapp.	configuration of the WebModeler webapp component	
webModeler.webapp.image	configuration of the webapp Docker image	
webModeler.webapp.image.repository	defines which image repository to use for the webapp Docker image	web-modeler-ee/modeler-webapp
webModeler.webapp.sidecars	can be used to attach extra containers to the modeler webapp deployment	[]
webModeler.webapp.initContainers	can be used to set up extra init containers for the application Pod	[]
webModeler.webapp.podAnnotations	can be used to define extra webapp pod annotations	{}
webModeler.webapp.podLabels	can be used to define extra webapp pod labels	{}
webModeler.webapp.env	can be used to set extra environment variables in each webapp container	[]
webModeler.webapp.command	can be used to override the default command provided by the container image, see <a href="https://kubernetes.io/docs/tasks/inject-data-application/define-command-argument-container/">https://kubernetes.io/docs/tasks/inject-data-application/define-command-argument-container/</a>	[]
webModeler.webapp.extraVolumes	can be used to define extra volumes for the webapp pods, useful for TLS and self-signed certificates	[]
webModeler.webapp.extraVolumeMounts	can be used to mount extra volumes for the webapp pods, useful for TLS and self-signed certificates	[]
webModeler.webapp.podSecurityContext	can be used to define the security options the webapp pod should be run with	
webModeler.webapp.podSecurityContext.runAsNonRoot		true
webModeler.webapp.podSecurityContext.fsGroup		1001
webModeler.webapp.containerSecurityContext	can be used to define the security options the webapp container should be run with	
webModeler.webapp.containerSecurityContext.privileged		false
webModeler.webapp.containerSecurityContext.readOnlyRootFilesystem		true
webModeler.webapp.containerSecurityContext.allowPrivilegeEscalation		false
webModeler.webapp.containerSecurityContext.runAsNonRoot		true
webModeler.webapp.containerSecurityContext.runAsUser		1001
webModeler.webapp.startupProbe	configuration of the webapp startup probe	
webModeler.webapp.startupProbe.enabled	if true, the startup probe will be enabled for the webapp container	false
webModeler.webapp.startupProbe.scheme	defines the startup probe schema used on calling the probePath	HTTP
webModeler.webapp.startupProbe.probePath	defines the HTTP endpoint used for the startup probe	/health/liveness
webModeler.webapp.startupProbe.initialDelaySeconds	defines the number of seconds after the container has started before the probe is initiated	15
webModeler.webapp.startupProbe.periodSeconds	defines how often the probe is executed	30
webModeler.webapp.startupProbe.successThreshold	defines how often the probe needs to succeed to be considered successful after having failed	1
webModeler.webapp.startupProbe.failureThreshold	defines when the probe is considered failed so the container will be restarted	5
webModeler.webapp.startupProbe.timeoutSeconds	defines the number of seconds after which the probe times out	1
webModeler.webapp.readinessProbe	configuration of the webapp readiness probe	
webModeler.webapp.readinessProbe.enabled	if true, the readiness probe will be enabled for the webapp container	true
webModeler.webapp.readinessProbe.scheme	defines the startup probe schema used on calling the probePath	HTTP
webModeler.webapp.readinessProbe.probePath	defines the HTTP endpoint used for the readiness probe	/health/readiness
webModeler.webapp.readinessProbe.initialDelaySeconds	defines the number of seconds after the container has started before the probe is initiated	15
webModeler.webapp.readinessProbe.periodSeconds	defines how often the probe is executed	30
webModeler.webapp.readinessProbe.successThreshold	defines how often the probe needs to succeed to be considered successful after having failed	1
webModeler.webapp.readinessProbe.failureThreshold	defines when the probe is considered failed so the Pod will be marked unready	5
webModeler.webapp.readinessProbe.timeoutSeconds	defines the number of seconds after which the probe times out	1
webModeler.webapp.livenessProbe	configuration of the webapp liveness probe	
webModeler.webapp.livenessProbe.enabled	if true, the liveness probe will be enabled for the webapp container	false

Name	Description	Value
<code>webModeler.webapp.livenessProbe.scheme</code>	defines the startup probe schema used on calling the probePath	HTTP
<code>webModeler.webapp.livenessProbe.probePath</code>	defines the HTTP endpoint used for the liveness probe	<code>/health/liveness</code>
<code>webModeler.webapp.livenessProbe.initialDelaySeconds</code>	defines the number of seconds after the container has started before the probe is initiated	15
<code>webModeler.webapp.livenessProbe.periodSeconds</code>	defines how often the probe is executed	30
<code>webModeler.webapp.livenessProbe.successThreshold</code>	defines how often the probe needs to succeed to be considered successful after having failed	1
<code>webModeler.webapp.livenessProbe.failureThreshold</code>	defines when the probe is considered failed so the container will be restarted	5
<code>webModeler.webapp.livenessProbe.timeoutSeconds</code>	defines the number of seconds after which the probe times out	1
<code>webModeler.webapp.metrics.prometheus</code>	Prometheus metrics endpoint	<code>/metrics</code>
<code>webModeler.webapp.nodeSelector</code>	can be used to select the nodes the webapp pods should run on	<code>{}</code>
<code>webModeler.webapp.tolerations</code>	can be used to define pod tolerations, see <a href="https://kubernetes.io/docs/concepts/scheduling-eviction/taint-and-toleration/">https://kubernetes.io/docs/concepts/scheduling-eviction/taint-and-toleration/</a>	<code>[]</code>
<code>webModeler.webapp.affinity</code>	can be used to define pod affinity or anti-affinity, see <a href="https://kubernetes.io/docs/concepts/scheduling-eviction/assign-pod-node/#affinity-and-anti-affinity">https://kubernetes.io/docs/concepts/scheduling-eviction/assign-pod-node/#affinity-and-anti-affinity</a>	<code>{}</code>
<code>webModeler.webapp.resources</code>	configuration of resource requests and limits for the container, see <a href="https://kubernetes.io/docs/concepts/configuration/manage-resources-containers/#requests-and-limits">https://kubernetes.io/docs/concepts/configuration/manage-resources-containers/#requests-and-limits</a>	
<code>webModeler.webapp.resources.requests.cpu</code>		400m
<code>webModeler.webapp.resources.requests.memory</code>		256Mi
<code>webModeler.webapp.resources.limits.cpu</code>		800m
<code>webModeler.webapp.resources.limits.memory</code>		512Mi
<code>webModeler.webapp.service</code>	configuration of the WebModeler webapp service	
<code>webModeler.webapp.service.annotations</code>	can be used to define annotations which will be applied to the service	<code>{}</code>
<code>webModeler.webapp.service.type</code>	defines the type of the service, see <a href="https://kubernetes.io/docs/concepts/services-networking/service/#publishing-services-service-types">https://kubernetes.io/docs/concepts/services-networking/service/#publishing-services-service-types</a>	ClusterIP
<code>webModeler.webapp.service.port</code>	defines the port of the service	80
<code>webModeler.webapp.service.managementPort</code>	defines the management port of the service	8071
<code>webModeler.webapp.configuration</code>	if specified, contents will be used as the application.yaml	<code>""</code>
<code>webModeler.webapp.extraConfiguration</code>	if specified, contents will be used for any extra configuration files such as log4j2.xml	<code>{}</code>

WebModeler - WebSockets Parameters

Name	Description	Value
<code>webModeler.websockets</code>	configuration of the WebModeler websockets component	
<code>webModeler.websockets.image</code>	configuration of the websockets Docker image	
<code>webModeler.websockets.image.repository</code>	defines which image repository to use for the websockets Docker image	<code>web-modeler-ee/modeler-websockets</code>
<code>webModeler.websockets.sidecars</code>	can be used to attach extra containers to the modeler websockets deployment	<code>[]</code>
<code>webModeler.websockets.initContainers</code>	can be used to set up extra init containers for the application Pod	<code>[]</code>
<code>webModeler.websockets.publicHost</code>	can be used to define the host on which the WebSockets server can be reached from the WebModeler client in the browser.	<code>localhost</code>
<code>webModeler.websockets.publicPort</code>	can be used to define the port number on which the WebSockets server can be reached from the WebModeler client in the browser.	8085
<code>webModeler.websockets.podAnnotations</code>	can be used to define extra websockets pod annotations	<code>{}</code>
<code>webModeler.websockets.podLabels</code>	can be used to define extra websockets pod labels	<code>{}</code>
<code>webModeler.websockets.env</code>	can be used to set extra environment variables in each websockets container	<code>[]</code>
<code>webModeler.websockets.command</code>	can be used to override the default command provided by the container image, see <a href="https://kubernetes.io/docs/tasks/inject-data-application/define-command-argument-container/">https://kubernetes.io/docs/tasks/inject-data-application/define-command-argument-container/</a>	<code>[]</code>
<code>webModeler.websockets.extraVolumes</code>	can be used to define extra volumes for the websockets pod; useful for logging to a file	<code>[]</code>
<code>webModeler.websockets.extraVolumeMounts</code>	can be used to mount extra volumes for the websockets pod; useful for logging to a file	<code>[]</code>
<code>webModeler.websockets.podSecurityContext</code>	can be used to define the security options the websockets pod should be run with	
<code>webModeler.websockets.podSecurityContext.runAsNonRoot</code>		true
<code>webModeler.websockets.podSecurityContext.fsGroup</code>		1001
<code>webModeler.websockets.containerSecurityContext</code>	can be used to define the security options the websockets container should be run with	
<code>webModeler.websockets.containerSecurityContext.privileged</code>		false
<code>webModeler.websockets.containerSecurityContext.readOnlyRootFilesystem</code>		true
<code>webModeler.websockets.containerSecurityContext.allowPrivilegeEscalation</code>		false

Name	Description	Value
<code>webModeler.websockets.containerSecurityContext.runAsNonRoot</code>		<code>true</code>
<code>webModeler.websockets.containerSecurityContext.runAsUser</code>		1001
<code>webModeler.websockets.startupProbe</code>	configuration of the websockets startup probe	
<code>webModeler.websockets.startupProbe.enabled</code>	if true, the startup probe will be enabled for the websockets container	<code>false</code>
<code>webModeler.websockets.startupProbe.initialDelaySeconds</code>	defines the number of seconds after the container has started before the probe is initiated	10
<code>webModeler.websockets.startupProbe.periodSeconds</code>	defines how often the probe is executed	30
<code>webModeler.websockets.startupProbe.successThreshold</code>	defines how often the probe needs to succeed to be considered successful after having failed	1
<code>webModeler.websockets.startupProbe.failureThreshold</code>	defines when the probe is considered failed so the container will be restarted	5
<code>webModeler.websockets.startupProbe.timeoutSeconds</code>	defines the number of seconds after which the probe times out	1
<code>webModeler.websockets.readinessProbe</code>	configuration of the websockets readiness probe	
<code>webModeler.websockets.readinessProbe.enabled</code>	if true, the readiness probe will be enabled for the websockets container	<code>true</code>
<code>webModeler.websockets.readinessProbe.initialDelaySeconds</code>	defines the number of seconds after the container has started before the probe is initiated	10
<code>webModeler.websockets.readinessProbe.periodSeconds</code>	defines how often the probe is executed	30
<code>webModeler.websockets.readinessProbe.successThreshold</code>	defines how often the probe needs to succeed to be considered successful after having failed	1
<code>webModeler.websockets.readinessProbe.failureThreshold</code>	defines when the probe is considered failed so the Pod will be marked unready	5
<code>webModeler.websockets.readinessProbe.timeoutSeconds</code>	defines the number of seconds after which the probe times out	1
<code>webModeler.websockets.livenessProbe</code>	configuration of the websockets liveness probe	
<code>webModeler.websockets.livenessProbe.enabled</code>	if true, the liveness probe will be enabled for the websockets container	<code>false</code>
<code>webModeler.websockets.livenessProbe.initialDelaySeconds</code>	defines the number of seconds after the container has started before the probe is initiated	10
<code>webModeler.websockets.livenessProbe.periodSeconds</code>	defines how often the probe is executed	30
<code>webModeler.websockets.livenessProbe.successThreshold</code>	defines how often the probe needs to succeed to be considered successful after having failed	1
<code>webModeler.websockets.livenessProbe.failureThreshold</code>	defines when the probe is considered failed so the container will be restarted	5
<code>webModeler.websockets.livenessProbe.timeoutSeconds</code>	defines the number of seconds after which the probe times out	1
<code>webModeler.websockets.nodeSelector</code>	can be used to select the nodes the websockets pods should run on	<code>{}</code>
<code>webModeler.websockets.tolerations</code>	can be used to define pod tolerations, see <a href="https://kubernetes.io/docs/concepts/scheduling-eviction/taint-and-toleration/">https://kubernetes.io/docs/concepts/scheduling-eviction/taint-and-toleration/</a>	<code>[]</code>
<code>webModeler.websockets.affinity</code>	can be used to define pod affinity or anti-affinity, see <a href="https://kubernetes.io/docs/concepts/scheduling-eviction/assign-pod-node/#affinity-and-anti-affinity">https://kubernetes.io/docs/concepts/scheduling-eviction/assign-pod-node/#affinity-and-anti-affinity</a>	<code>{}</code>
<code>webModeler.websockets.resources</code>	configuration of resource requests and limits for the container, see <a href="https://kubernetes.io/docs/concepts/configuration/manage-resources-containers/#requests-and-limits">https://kubernetes.io/docs/concepts/configuration/manage-resources-containers/#requests-and-limits</a>	
<code>webModeler.websockets.resources.requests.cpu</code>		100m
<code>webModeler.websockets.resources.requests.memory</code>		64Mi
<code>webModeler.websockets.resources.limits.cpu</code>		200m
<code>webModeler.websockets.resources.limits.memory</code>		128Mi
<code>webModeler.websockets.service</code>	configuration of the WebModeler websockets service	
<code>webModeler.websockets.service.annotations</code>	can be used to define annotations which will be applied to the service	<code>{}</code>
<code>webModeler.websockets.service.type</code>	defines the type of the service, see <a href="https://kubernetes.io/docs/concepts/services-networking/service/#publishing-services-service-types">https://kubernetes.io/docs/concepts/services-networking/service/#publishing-services-service-types</a>	<code>ClusterIP</code>
<code>webModeler.websockets.service.port</code>	defines the port of the service	80
<code>webModeler.websockets.configuration</code>	if specified, contents will be used as the application.yaml	<code>""</code>
<code>webModeler.websockets.extraConfiguration</code>	if specified, contents will be used for any extra configuration files such as log4j2.xml	<code>{}</code>
<code>webModeler.serviceAccount</code>	configuration for the service account the WebModeler pods are assigned to	
<code>webModeler.serviceAccount.enabled</code>	if true, enables the WebModeler service account	<code>true</code>
<code>webModeler.serviceAccount.name</code>	can be used to set the name of the WebModeler service account	<code>""</code>
<code>webModeler.serviceAccount.annotations</code>	can be used to set the annotations of the WebModeler service account	<code>{}</code>
<code>webModeler.serviceAccount.automountServiceAccountToken</code>	can be used to control whether the service account token should be automatically mounted	<code>false</code>
<code>webModeler.ingress.enabled</code>	if true, an Ingress resource will be deployed with the WebModeler deployment. Only useful if an Ingress controller like NGINX is available.	<code>false</code>
<code>webModeler.ingress.className</code>	defines the class or configuration of ingress which should be used by the controller	<code>nginx</code>
<code>webModeler.ingress.annotations</code>	defines the ingress related annotations, consumed mostly by the ingress controller	<code>{}</code>
<code>webModeler.ingress.webapp</code>	configuration of the webapp ingress	
<code>webModeler.ingress.webapp.host</code>	defines the host of the ingress rule, see <a href="https://kubernetes.io/docs/concepts/services-networking/ingress/#ingress-rules">https://kubernetes.io/docs/concepts/services-networking/ingress/#ingress-rules</a> ; this is the host name on which the WebModeler web application will be available	<code>""</code>

Name	Description	Value
webModeler.ingress.webapp.pathType	can be used to define the Ingress path type. <a href="https://kubernetes.io/docs/concepts/services-networking/ingress/#path-types">https://kubernetes.io/docs/concepts/services-networking/ingress/#path-types</a>	Prefix
webModeler.ingress.webapp.tls	configuration for TLS on the ingress resource, see <a href="https://kubernetes.io/docs/concepts/services-networking/ingress/#tls">https://kubernetes.io/docs/concepts/services-networking/ingress/#tls</a>	
webModeler.ingress.webapp.tls.enabled	if true, TLS will be configured on the ingress resource	false
webModeler.ingress.webapp.tls.secretName	defines the secret name which contains the TLS private key and certificate	camunda-platform-webmodeler-webapp
webModeler.ingress.websockets	configuration of the websockets ingress	
webModeler.ingress.websockets.host	defines the host of the ingress rule, see <a href="https://kubernetes.io/docs/concepts/services-networking/ingress/#ingress-rules">https://kubernetes.io/docs/concepts/services-networking/ingress/#ingress-rules</a> ; this is the host name the WebModeler client in the browser will use to connect to the WebSockets server	""
webModeler.ingress.websockets.pathType	can be used to define the Ingress path type. <a href="https://kubernetes.io/docs/concepts/services-networking/ingress/#path-types">https://kubernetes.io/docs/concepts/services-networking/ingress/#path-types</a>	Prefix
webModeler.ingress.websockets.tls	configuration for TLS on the ingress resource, see <a href="https://kubernetes.io/docs/concepts/services-networking/ingress/#tls">https://kubernetes.io/docs/concepts/services-networking/ingress/#tls</a>	
webModeler.ingress.websockets.tls.enabled	if true, TLS will be configured on the ingress resource	false
webModeler.ingress.websockets.tls.secretName	defines the secret name which contains the TLS private key and certificate	camunda-platform-webmodeler-websockets

WebModeler - PostgreSQL Parameters

Name	Description	Value
postgresql	configuration for the postgresql dependency chart used by WebModeler. See the chart documentation <a href="https://github.com/bitnami/charts/tree/master/bitnami/postgresql#parameters">https://github.com/bitnami/charts/tree/master/bitnami/postgresql#parameters</a> for more details.	
postgresql.enabled	if true, a PostgreSQL database will be deployed as part of the Helm release by using the dependency chart	false
postgresql.nameOverride	defines the name of the Postgres resources (names will be prefixed with the release name), see <a href="https://github.com/bitnami/charts/tree/main/bitnami/postgresql#common-parameters">https://github.com/bitnami/charts/tree/main/bitnami/postgresql#common-parameters</a>	postgresql-web-modeler
postgresql.auth	configuration of the database authentication	
postgresql.auth.username	defines the name of the database user to be created for WebModeler	web-modeler
postgresql.auth.password	defines the database user’s password; a random password will be generated if left empty	""
postgresql.auth.database	defines the name of the database to be created for WebModeler	web-modeler
postgresql.primary.extraVolumes[0].name		tmp
postgresql.primary.extraVolumes[0].emptyDir		{}
postgresql.primary.extraVolumes[1].name		config
postgresql.primary.extraVolumes[1].emptyDir		{}
postgresql.primary.extraVolumes[2].name		postgresql-tmp
postgresql.primary.extraVolumes[2].emptyDir		{}
postgresql.primary.extraVolumeMounts[0].mountPath		/tmp
postgresql.primary.extraVolumeMounts[0].name		tmp
postgresql.primary.extraVolumeMounts[1].mountPath		/opt/bitnami/postgresql/conf
postgresql.primary.extraVolumeMounts[1].name		config
postgresql.primary.extraVolumeMounts[2].mountPath		/opt/bitnami/postgresql/tmp
postgresql.primary.extraVolumeMounts[2].name		postgresql-tmp
postgresql.primary.containerSecurityContext.enabled		true
postgresql.primary.containerSecurityContext.privileged		false
postgresql.primary.containerSecurityContext.readOnlyRootFilesystem		true
postgresql.primary.containerSecurityContext.allowPrivilegeEscalation		false
postgresql.primary.containerSecurityContext.runAsNonRoot		true
postgresql.primary.containerSecurityContext.runAsUser		1001
postgresql.primary.containerSecurityContext.capabilities.drop		["ALL"]
postgresql.primary.containerSecurityContext.seccompProfile.type		RuntimeDefault
postgresql.primary.podSecurityContext.enabled		true
postgresql.primary.podSecurityContext.runAsNonRoot		true
postgresql.primary.podSecurityContext.fsGroup		1001

Connectors Parameters

Name	Description	Value
connectors	configuration for the Connectors.	
connectors.enabled	if true, the Connectors deployment and its related resources are deployed via a helm release	true
connectors.inbound	Switch for inbound mode (e.g., for webhook or polling)	
connectors.inbound.mode	acceptable values: disabled, credentials, or oauth	oauth
connectors.inbound.auth	configuration of the credentials authentication.	
connectors.inbound.auth.existingSecret	can be used to configure Secret name that contains Operate password (if inbound mode is credentials)	""
connectors.image	configuration to configure the Connectors image specifics	
connectors.image.registry	can be used to set container image registry.	""
connectors.image.repository	defines which image repository to use	camunda/connectors-bundle
connectors.image.tag	can be set to overwrite the global tag, which should be used in that chart	8.5.0
connectors.image.pullSecrets	can be used to configure image pull secrets <a href="https://kubernetes.io/docs/concepts/containers/images/#specifying-imagepullsecrets-on-a-pod">https://kubernetes.io/docs/concepts/containers/images/#specifying-imagepullsecrets-on-a-pod</a>	[]
connectors.sidecars	can be used to attach extra containers to the connectors deployment	[]
connectors.initContainers	can be used to set up extra init containers for the application Pod	[]
connectors.replicas	number of Connectors replicas	1
connectors.contextPath	can be used to make Connectors web application works on a custom sub-path. This is mainly used to run Camunda web applications under a single domain.	""
connectors.podAnnotations	can be used to define extra Connectors pod annotations	{}
connectors.podLabels	can be used to define extra Connectors pod labels	{}
connectors.logging	configuration for the Connectors logging. This template will be directly included in the Operate configuration YAML file	
connectors.logging.level.io.camunda.connector		ERROR
connectors.service	configuration to configure the Connectors service.	
connectors.service.annotations	can be used to define annotations, which will be applied to the Connectors service	{}
connectors.service.type	defines the type of the service <a href="https://kubernetes.io/docs/concepts/services-networking/service/#publishing-services-service-types">https://kubernetes.io/docs/concepts/services-networking/service/#publishing-services-service-types</a>	ClusterIP
connectors.service.serverPort	defines the port number where the Connector web application will be available	8080
connectors.service.serverName	defines the port name where the Connector web application will be available	http
connectors.resources	configuration to set request and limit configuration for the container <a href="https://kubernetes.io/docs/concepts/configuration/manage-resources-containers/#requests-and-limits">https://kubernetes.io/docs/concepts/configuration/manage-resources-containers/#requests-and-limits</a>	
connectors.resources.requests.cpu		1
connectors.resources.requests.memory		1Gi
connectors.resources.limits.cpu		2
connectors.resources.limits.memory		2Gi
connectors.env	can be used to set extra environment variables in each Connector container	[]
connectors.command	can be used to override the default command provided by the container image. See <a href="https://kubernetes.io/docs/tasks/inject-data-application/define-command-argument-container/">https://kubernetes.io/docs/tasks/inject-data-application/define-command-argument-container/</a>	[]
connectors.extraVolumes	can be used to define extra volumes for the Connectors pods, useful for TLS and self-signed certificates	[]
connectors.extraVolumeMounts	can be used to mount extra volumes for the Connectors pods, useful for TLS and self-signed certificates	[]
connectors.startupProbe	configuration	
connectors.startupProbe.enabled	if true, the startup probe is enabled in app container	false
connectors.startupProbe.scheme	defines the startup probe scheme used on calling the probePath	HTTP
connectors.startupProbe.probePath	defines the startup probe route used on the app	/actuator/health/readiness
connectors.startupProbe.initialDelaySeconds	defines the number of seconds after the container has started before	30
connectors.startupProbe.periodSeconds	defines how often the probe is executed	30
connectors.startupProbe.successThreshold	defines how often it needs to be true to be marked as ready, after failure	1
connectors.startupProbe.failureThreshold	defines when the probe is considered as failed so the Pod will be marked Unready	5
connectors.startupProbe.timeoutSeconds	defines the seconds after the probe times out	1
connectors.readinessProbe	configuration	
connectors.readinessProbe.enabled	if true, the readiness probe is enabled in app container	true
connectors.readinessProbe.scheme	defines the startup probe scheme used on calling the probePath	HTTP
connectors.readinessProbe.probePath	defines the readiness probe route used on the app	/actuator/health/readiness
connectors.readinessProbe.initialDelaySeconds	defines the number of seconds after the container has started before the probe is initiated.	30
connectors.readinessProbe.periodSeconds	defines how often the probe is executed	30
connectors.readinessProbe.successThreshold	defines how often it needs to be true to be marked as ready, after failure	1
connectors.readinessProbe.failureThreshold	defines when the probe is considered as failed so the Pod will be marked Unready	5

Name	Description	Value
connectors.readinessProbe.timeoutSeconds	defines the seconds after the probe times out	1
connectors.livenessProbe	configuration	
connectors.livenessProbe.enabled	if true, the liveness probe is enabled in app container	false
connectors.livenessProbe.scheme	defines the startup probe scheme used on calling the probePath	HTTP
connectors.livenessProbe.probePath	defines the liveness probe route used on the app	/actuator/health/liveness
connectors.livenessProbe.initialDelaySeconds	defines the number of seconds after the container has started before	30
connectors.livenessProbe.initialDelaySeconds	the probe is initiated.	30
connectors.livenessProbe.periodSeconds	defines how often the probe is executed	30
connectors.livenessProbe.successThreshold	defines how often it needs to be true to be considered successful after having failed	1
connectors.livenessProbe.failureThreshold	defines when the probe is considered as failed so the container will be restarted	5
connectors.livenessProbe.timeoutSeconds	defines the seconds after the probe times out	1
connectors.metrics.prometheus	Prometheus metrics endpoint	/actuator/prometheus
connectors.serviceAccount	configuration for the service account where the Connectors pods are assigned to	
connectors.serviceAccount.enabled	if true, enables the Connectors service account	false
connectors.serviceAccount.name	can be used to set the name of the Connectors service account	""
connectors.serviceAccount.annotations	can be used to set the annotations of the Operate service account	{}
connectors.serviceAccount.automountServiceAccountToken	can be used to control whether the service account token should be automatically mounted	false
connectors.ingress.enabled	if true, an ingress resource is deployed with the Connectors deployment. Only useful if an ingress controller is available, like nginx.	false
connectors.ingress.className	defines the class or configuration of ingress which should be used by the controller	nginx
connectors.ingress.annotations	defines the ingress related annotations, consumed mostly by the ingress controller	{}
connectors.ingress.path	defines the path which is associated with the Connectors service and port	/
connectors.ingress.pathType	https://kubernetes.io/docs/concepts/services-networking/ingress/#ingress-rules can be used to define the Ingress path type.	Prefix
connectors.ingress.host	https://kubernetes.io/docs/concepts/services-networking/ingress/#path-types can be used to define the host of the ingress rule.	""
connectors.ingress.tls	https://kubernetes.io/docs/concepts/services-networking/ingress/#ingress-rules configuration for tls on the ingress resource	
connectors.ingress.tls.enabled	https://kubernetes.io/docs/concepts/services-networking/ingress/#tls if true, then tls is configured on the ingress resource. If enabled the Ingress.host need to be defined.	false
connectors.ingress.tls.secretName	defines the secret name which contains the TLS private key and certificate	camunda-platform-connectors
connectors.podSecurityContext	defines the security options the Connectors pod should be run with	
connectors.podSecurityContext.runAsNonRoot	run as non root	true
connectors.podSecurityContext.fsGroup		1001
connectors.containerSecurityContext	defines the security options the Connectors container should be run with	
connectors.containerSecurityContext.privileged		false
connectors.containerSecurityContext.readOnlyRootFilesystem		true
connectors.containerSecurityContext.allowPrivilegeEscalation		false
connectors.containerSecurityContext.runAsNonRoot		true
connectors.containerSecurityContext.runAsUser		1001
connectors.nodeSelector	can be used to define on which nodes the Connectors pods should run	{}
connectors.tolerations	can be used to define pod toleration's https://kubernetes.io/docs/concepts/scheduling-eviction/taint-and-toleration/	[]
connectors.affinity	can be used to define pod affinity or anti-affinity https://kubernetes.io/docs/concepts/scheduling-eviction/assign-pod-node/#affinity-and-anti-affinity	{}
connectors.configuration	if specified, contents will be used as the application.yaml	""
connectors.extraConfiguration	if specified, contents will be used for any extra configuration files such as the log4j2.xml	{}

#### Elasticsearch Parameters

Name	Description	Value
elasticsearch		
elasticsearch.enabled		true
elasticsearch.image.repository		bitnami/elasticsearch
elasticsearch.image.tag		8.12.2
elasticsearch.master.containerSecurityContext.readOnlyRootFilesystem		true
elasticsearch.master.masterOnly		false
elasticsearch.master.heapSize		1024m

Name	Description	Value
elasticsearch.master.persistence.size		64Gi
elasticsearch.master.resources.requests.cpu	cpu request	1
elasticsearch.master.resources.requests.memory	request	2Gi
elasticsearch.master.resources.limits.cpu	cpu limit	2
elasticsearch.master.resources.limits.memory	memory limit	2Gi
elasticsearch.master.extraEnvVars[0].name	env	ELASTICSEARCH_ENABLE_REST_TLS
elasticsearch.master.extraEnvVars[0].value	env value	false
elasticsearch.sysctlImage.enabled		true
elasticsearch.data.replicaCount		0
elasticsearch.coordinating.replicaCount		0
elasticsearch.ingest.enabled		false

Prometheus Parameters

Name	Description	Value
PrometheusServiceMonitor	configuration to configure a prometheus service monitor	
prometheusServiceMonitor.enabled	if true then a service monitor will be deployed, which allows an installed prometheus controller to scrape metrics from the deployed pods	false
promotheuServiceMonitor.labels	can be set to configure extra labels, which will be added to the servicemonitor and can be used on the prometheus controller for selecting the servicemonitors	
prometheusServiceMonitor.labels.release		metrics
prometheusServiceMonitor.scrapeInterval	can be set to configure the interval at which metrics should be scraped	10s